

Research note

Michael J. Lopez*

Analyzing the National Football League is challenging, but player tracking data is here to help

Abstract: Most historical National Football League (NFL) analysis, both mainstream and academic, has relied on public, play-level data to generate team and player comparisons. Given the number of oft omitted variables that impact on-field results, such as play call, game situation, and opponent strength, findings tend to be more anecdotal than actionable. With the release of player tracking data, however, analysts can better ask and answer questions to isolate skill and strategy. In this article, we highlight the limitations of traditional analyses, and use a decades-old punching bag for analysts – fourth-down strategy – as a microcosm for why tracking data is needed. Specifically, we assert that, in absence of a previously unmeasured variable, past findings supporting an aggressive fourth down strategy may have been overstated. Next, we synthesize recent work that comprises this special Journal of Quantitative Analysis in Sports issue into player tracking data in football. Finally, we conclude with some best practices and limitations regarding usage of this data. The release of player tracking data marks a transition for the league and its’ analysts, and we hope this issue helps guide innovation in football analytics for years to come.

Keywords: National Football League; player tracking; causal inference; fourth downs

1 The problem with football data

Across the physical and social sciences, the gold standard for identifying causes and effects of certain behaviors, therapies, or interventions is the randomized experiment. Randomization is attractive because subjects that receive one treatment are, in expectation, comparable to those that receive another treatment. When examining an outcome of interest in a randomized trial, one can be confident that there are no underlying and unmeasurable differences (e.g, confounding variables) that would

*Corresponding author: Michael J. Lopez, National Football League

be responsible for causing said results. Succinctly – randomized designs ensure that apples are compared to apples.

Analyzing football data is like that, except the opposite. That is, all virtually all football data is observational, which means that any study of athlete or team behavior is potentially confounded by other variables linked to the game or player. Because of this, it is often quite difficult to rule out whether extraneous factors related to players and games are responsible for findings. In football, we rarely get to compare apples to apples.

Consider the well-established argument in NFL analytics circles that teams should pass more and run less (Kovash and Levitt, 2009; Hermsmeyer, 2018). Unlike what would happen in randomized designs, teams don't flip a coin to decide if they'll attempt a pass. Instead, coaches call a pass play based on, among other factors, quarterback skill, game situation, and opposing defense, many of which are traits that likely differ from settings that call for run plays. Certain drivers of play choice, including down, distance needed for a first down, and score differential, are known prior to a play. But several other variables, such as quarterback health, defensive formation and personnel, and pre-snap movement, are both (i) likewise linked to whether or not a team attempts a pass, and (ii) difficult to quantify.

Not only are there several factors that go into NFL play-calling, but many cannot be measured using traditional data. An artifact of the NFL's historical data collection is that the hundreds of player movements and decisions in a play are reduced to one observation, one row in a data set. Worse, at least in terms of public data, most of the 22 players on the field at a given time aren't even recorded as being there (Schatz, 2005). At around 160 plays per game, it is feasible to look at play-level outcomes such as win probability, expected points, run/pass strategy, and fourth down behavior, while conditioning on what we know about plays before they happen. But both before and within a play, events in football are reactionary. Substitutions by one team lead to changes in personnel by the other. Formations by the defense lead to audibles by the offense. Motion from the slot back spurs a new defensive coverage scheme. And even after the play, hip placement of defensive backs creates new cuts from wide receivers and a lineman's first step signals an opposing linebacker's read. Under this complex interplay, nothing is as simple as whether or not a team's coach called a passing play, which makes trying to untangle the marginal effect of passing versus running exceedingly difficult. Thus, even when conditioning on several variables, such as Hermsmeyer (2018), which includes number of defenders near the line-of-scrimmage and other play and game-level characteristics, we are likely not accounting for something else.

Why do these differences matter? The answer lies in how we leverage analytical insight to impact the game. It means one thing to find that passing results in better outcomes (yards per play, expected yards per play, win probability added) than run

plays. Analysis can be improved by conditioning on traits such as down, distance, and formation. But it means substantially more to find that among teams that ran the ball, they'd have been better off had they passed more often. This is a tricky but important distinction. This second claim is a causal one, and cannot be made with certitude unless we are able to condition on all relevant football variables. Researchers are unable to condition on variables we do not have, which jeopardizes any attempt to establish causality in football.

Beginning in 2016, those 160 observations per game turned into roughly 300,000. That season, each player was equipped with radio-frequency identification (RFID) chips in each of his shoulder pads, observations that provide the location of each player at roughly 10 frames-per-second, wherever he goes on the field. From player movement, it is straightforward to calculate speed, angle, and acceleration. These data are termed the NFL's "Next Gen Stats" (see <https://nextgenstats.nfl.com/> for more insight and summary metrics). For the first time, analysts can dig beyond play-level analyses to better understand the game. Variables that used to only show up in scouting reports – ones such as route type, running back space, or defensive back coverage – can now be estimated directly from data. Soon, analysts will be breaking down a game before their coaches have even watched the film.

Given the novelty of player tracking data, much remains unknown about the questions that can be answered from it. But for statisticians, the goal should be clear – how can we leverage this data to ensure we are comparing apples to apples?

2 The importance of unmeasured confounding: a unique example

Over the last several years, the lowest hanging fruit among NFL analytics enthusiasts has been team behavior on fourth down. Authors as far back as Carter and Machol (1978), and including more recent work from Romer (2006), Burke et al. (2013), and the author of this manuscript (Yam and Lopez (2019)), have argued that teams are too passive on fourth down. Roughly, it has been argued that a more aggressive strategy is worth 0.4 wins per year (Romer, 2006; Yam and Lopez, 2019). In a limited 16-game season, that fractional gain takes on an added importance.

Traditional fourth-down analyses has compared play outcomes (such as success rates, expected points, or win probability) before and after potential conversion attempts. The crux of these works requires some level of extrapolation regarding the outcomes for teams that did not go for it, if they were instead to have gone for it. For example, Romer (2006) used success rate, Burke and Quealy (2013) used

expected points, and Yam and Lopez (2019) used win probability, all to imply that teams that did not go for it would have been better off going for it.

Each fourth-down analysis framework mentioned above assumed some level of equipoise between teams that did not go for it and those that did. Yam and Lopez (2019) went as far as using propensity-score based matching techniques to ensure that teams that did not go for it were compared to similar teams that did go for it. But even the most novel of matching techniques designed for observational data cannot get around the limitation of unmeasured confounding. That is, what if there was an unmeasured characteristic of teams that went for it that fundamentally differed from those that did not go for it, one that simultaneously impacted both coaches' decisions and play outcomes?

Turns out, there is one – the actual distance needed for a first down.

2.1 How actual distance impacts fourth-down strategy

In the gathering and disseminating of play-by-play data, NFL game-day assistants assign an integer value to each distance needed for a first down. All distances between 0 and 2 yards (not inclusive) are supposed to be listed as 4th-and-1's, with successive buckets consisting of one-yard intervals. So, both 4th-and-0.1 inches and 4th-and-71.9 inches are listed as 4th-and-1's. Likewise, distances of 4th-and-72 inches to 4th-and-107.9 inches are intended to be listed as 4th-and-2's.

Figure 1 shows two pairs of density plots containing the distributions of actual yards needed for a first down, comparing the distance on the RFID chip embedded in the football (prior to the ball snap) to the RFID chip in the sideline chain.¹ The chart is split by the integer categorization in the NFL play-by-play data (either a 4th-and-1 or a 4th-and-2).

Teams that went for it on 4th-and-1 did so from an average of 0.78 yards away from the line to gain; teams that did not go for it, meanwhile, were 0.99 yards away. On 4th-and-2 plays, teams that went for it did so from an average of 2.02 yards away, while those that did not averaged 2.10 yards away.²

¹ On goal-to-go plays, the actual distance needed for the first down is the difference between the ball RFID chip and the goal-line.

² We also used a second data wrangling strategy, where we identified the distance needed for a first down by using ten-yards from the original spot where each ball was marked on first down (with an exception taken for goal-line plays). On fourth-down plays, the differences between teams that went for it and teams that did not go for it were slightly larger than the ones shown here.

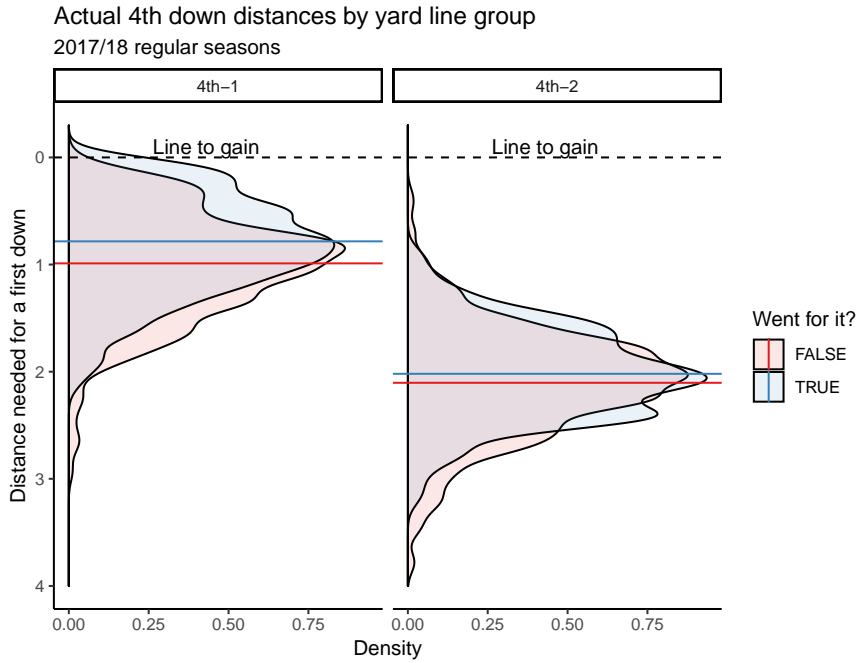


Fig. 1: Exact distances needed for a first down, split by teams that went for it (yes/no) and NFL play-by-play yard line categorization (4th-and-1, 4th-and-2). Teams that went for it did so from shorter distances within each yard line group.

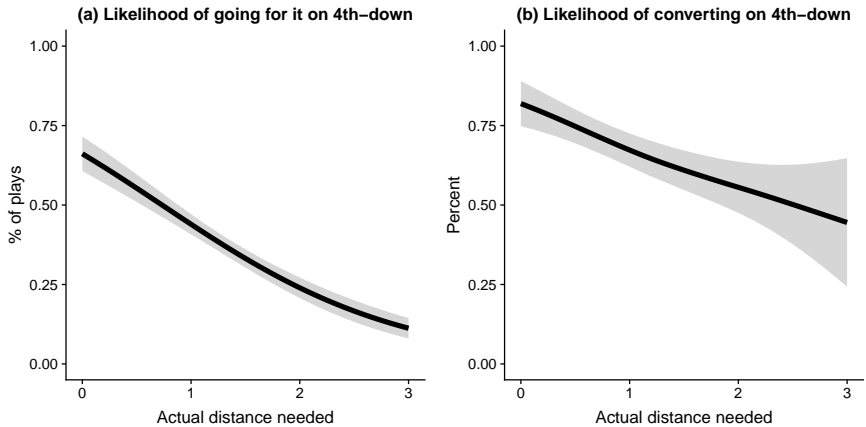


Fig. 2: Likelihood of going for it (left side) on 4th-down, and likelihood of a successful conversion (right side) on 4th-down. Each line shows the estimates from a generalized model of outcome (either going for it, or of successfully going for it) given the actual distance needed for a first down. Teams with shorter distances are more likely to go for it and to convert.

The actual distance needed for a first down yardage impacts both the attempt rate (among all teams) and the success rate (among teams that went for it), as shown in Figure 2. The left side of Figure 2 provides estimates from a generalized additive model (GAM) of attempt rate (Did teams go for it (Y/N), conditional on the actual distance needed for a first down), fit using 2017 and 2018 fourth down plays. Teams in 4th-and-inches went for it about 70% of the time, while teams in *long* 4th-and-1 situations went for it about 30% of the time. The right side of Figure 2 highlights how the rate of success varied based on actual distance needed for a first down, using estimates from a separate generalized additive model (Did teams gain a first down (Y/N), conditional on the actual distance needed). On 4th-and-inches, teams converted roughly 82% of the time, while in *long* 4th-and-1 situations, they converted around 55% of the time.

What does this imply?

Because they had further to go for a first down, teams that did not go for it did not have the same chances for success as the teams that did go for it. Thus, findings that inferred success rates, expected points, or win probability outcomes from teams that went for it almost assuredly overestimated the benefit of going for it on fourth down. For years, we failed to compare apples to apples.

Interestingly, although the GAM's in Figure 2 allow for a non-linear relationship between distance and each outcome, each curve looks somewhat linear. If we instead fit a line using the integer categorization of distance, we would see roughly the same

figure. That is, even though we were able to use more precise data, our results still matched how we could have predicted coaches to behave using only static data.

Replicating the approach of Yam and Lopez (2019) can help approximate what this means as far as the value of 4th-down aggressiveness. While those authors used data from prior to the 2016 season, we use the 2017 and 2018 seasons only, with an identical approach and code provided at <https://github.com/statsbylopez/nfl-fourth-down/tree/master/Code>. We replicate under two scenarios. First, we use the play-by-play distance category to check results over the past two seasons. Second, we use the actual distance that the offense needed for a first down in place of the interger distance. As in Yam and Lopez (2019), we only use plays in the “go for it” range of the 4th-down-bot (Burke et al., 2013).

Using the play-by-play distance category (4th-and-1, 4th-and-2, etc), we find that an aggressive fourth-down strategy would have been worth, in expectation, an extra 0.27 wins per-team per-year during the 2017 and 2018 seasons. Among teams that did not go for, we find an estimated 3.3% difference in win probability added per-play had they instead gone for it (95% CI, 1.9% to 4.7%). This is somewhat in line with the original findings, with a slight drop likely driven by recent upticks in team aggressiveness on fourth down (Stuart, 2017).

When using the actual distance needed for a first down (instead of the integer distance), the benefit of an aggressive fourth-down strategy drops to an extra 0.14 wins per year. Alternatively, the average difference in win probability added per-play is estimated at 2.0% (95% CI, 0.0% to 3.9%). For both per-play win probability and per-team benefit in terms of wins-added per-year, roughly 45% of the benefit to an aggressive fourth down strategy is negated when accounting for actual distance needed for a first down, a previously unmeasured confounding variable.

Though this attenuation of the effect size matches our intuition that going for it on fourth down is generally less advisable from longer distances, the primary lesson here extends beyond 4th-down strategy. Indeed, there could be other reasons why the benefit of aggressive approach remains underestimated (Romer, 2006). Instead, we aim to highlight the surprising insight that tracking data can provide. Repeated analyses of something in the game of football, in this case, fourth downs, had told us that coaches should “stop being stupid” with how they were acting (Cheema, 2019). But there was more to the story, in this case something intrinsically different about teams that went for it and teams that did not, which was responsible for at least part of the original findings.

Indeed, in football, there’s almost always more to the story.

3 Crowdsourcing insight into player tracking data

Papers accepted into this special issue of the Journal of Quantitative Analysis in Sports highlight the multitude of ideas where such apples-to-apples comparisons are now, more than ever, feasible.

Each of these papers stem from the National Football League's Big Data Bowl, an event originating from the league office designed to crowdsource public insight into tracking data, with the hope of inspiring analytically-inclined fans, curating ideas for team staffers, and spurring data-driven innovation in football.

****use examples here (this part will be based on which articles are accepted)****

4 Conclusion and next steps

This manuscript highlights several use cases of NFL tracking data, including both old and new research questions. In particular, we return to one of the league's oldest findings – that teams should be more aggressive on fourth down – to suggest that previous work may have overestimated the effect of going for it. Specifically, given the relative distance needed for a first down, teams that went for it on fourth down tended to do so from shorter distances, even when conditioning on the play-by-play yardage category. More generally, we summarize how the articles in this JQAS special issue on player tracking data will help shape the future of NFL analytics work.

Although we explicate on the value of player tracking insight in the NFL, it is important to acknowledge that this data is not a panacea for all football problems. Given the complexity of the game, there will always be fundamental football questions that data alone cannot precisely answer. Additionally, player tracking data is more arduous to analyze when compared to traditional play-by-play data; anecdotally, nearly every entrant to the league's Big Data Bowl wished they could have had more time to refine their work. To wit, here is a list of best practices and caveats for working with tracking data in the NFL.

1. Tracking data contains the x and y coordinates for each player and the football, collected at roughly 10 frames-per-second. Locational information is provided by signals sent from radio-frequency identification (RFID) chips that are placed inside each player's shoulder pads and inside the football. Speed, orientation, and distance traveled are straightforward to calculate using the tracking information, and are provided by the NFL's Next Gen Stats group. The typical game can contain

anywhere from 250,000 to 350,000 rows of data (1 row for each player on the field, on each play, at each time stamp) on which actual game action is occurring. Players are also tracked before and after plays (this information was not provided as part of the Big Data Bowl, and is generally seen as less pertinent). The z coordinate is not measured (e.g, height of the player, or height of the ball), nor can the precise location of helmets, arms, and legs be verified or easily estimated.

2. The field coordinates are fixed at each NFL stadium, as shown in Figure 3. From left to right, the length of the field spans from $x = 0$ to $x = 120$ (units are in yards), while the width of the field spans $y = 0$ to $y = 160/3$. Often, the first step in any analysis of tracking data is to ensure offensive teams are moving in the same direction. This requires flipping roughly half of a game's offensive plays from one direction to the other, while creating new x (subtracted from 120) and y (subtracted from $160/3$) coordinates. Additionally, standardizing by the play's line-of-scrimmage may be warranted.

3. Several play-specific traits remain unknown even when looking at tracking data. These include the initial play call, if the quarterback or coach called an audible, how the defense would have lined up if the offense used a different formation, if a wide receiver ran the correct route, etc. Each of these variables may be pertinent to more precisely estimate, as an example, the value of passing versus running. The absence of important play-level qualities highlights the need for analysts to work directly and cohesively with football experts in order to maximize the value of tracking data.

4. Some specifics about the data that researchers may want to be aware of. First, given updates to the RFID tags prior to the start of each season, small differences in speed measurements may exist from one year to the next. Additionally, the coordinates on the football are considered to be slightly less reliable than the coordinates on the players. Next, while analyzing maximum speed for players is often an easy-to-understand step, researchers should be wary that occasionally this maximum speed is reached while (or immediately after) a player is hit or tackled by an opponent. That said, tracking data is considered quite dependable; according to the Next Gen Stats group, location information is accurate to within ± 12 inches, and reliable data has been collected on 99.999% of the entirety of players and games over the last two seasons.

For years, data-driven innovation in the football was limited, and the NFL was, rightly or not, perceived to be trailing other leagues in terms of how teams used analytics. But insight lagged, in part, because so too did data. Behind player tracking insight, such excuses are no longer valid. The NFL's new data cannot tell us where exactly to look for insight, but it will allow us to both create new stories and to make old ones more complete.

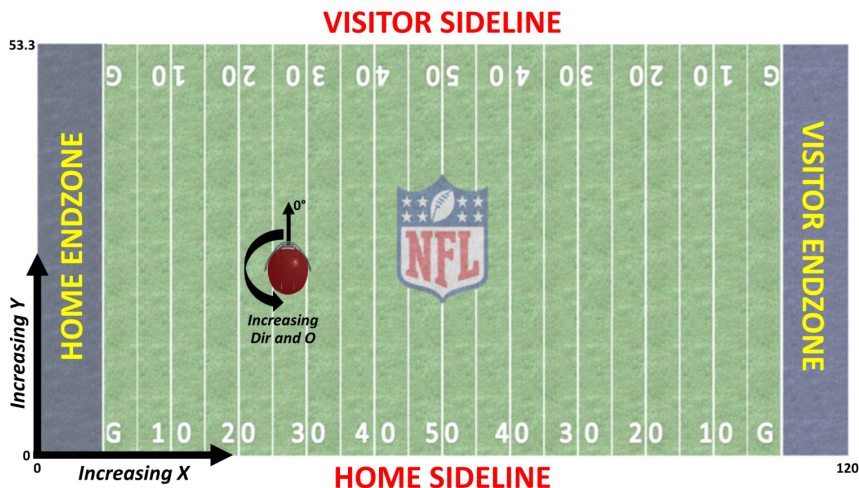


Fig. 3: Coordinates for working with player tracking data. Each stadium is equipped such that the home and visiting team end zones are fixed throughout the game. However, the end zones that teams defend in a game are not determined until the start of each half, and those directions change at the conclusion of the first and third quarters.

References

- B. Burke and K. Quealy. How coaches and the nyt 4th down bot compare. <http://www.nytimes.com/newsgraphics/2013/11/28/fourth-downs/post.html>, 2013. Accessed June 21, 2017.
- B. Burke, S. Carter, J. Daniel, T. Giratikanon, and K. Quealy. 4th down: When to go for it and why. https://www.nytimes.com/2014/09/05/upshot/4th-down-when-to-go-for-it-and-why.html?_r=0, 2013. Accessed March 20, 2017.
- V. Carter and R. E. Machol. Note—optimal strategies on fourth down. *Management Science*, 24(16):1758–1762, 1978.
- A. Cheema. Nfl coaches should stop being stupid on fourth downs. <https://www.thespax.com/nfl/message-to-nfl-coaches-stop-being-stupid-on-fourth-downs/>, 2019. Accessed July 20, 2019.
- J. Hermesmyer. For a passing league, the nfl still doesn't pass enough. <https://fivethirtyeight.com/features/for-a-passing-league-the-nfl-still-doesnt-pass-enough/>, 2018. Accessed June 27, 2019.
- K. Kovash and S. D. Levitt. Professionals do not play minimax: evidence from major league baseball and the national football league. Technical report, National Bureau of Economic Research, 2009.
- D. Romer. Do firms maximize? evidence from professional football. *Journal of Political Economy*, 114(2):340–365, 2006.
- A. Schatz. Football's hilbert problems. *Journal of Quantitative Analysis in Sports*, 1(1), 2005.

- C. Stuart. Nfl coaches are making better decisions. <https://slate.com/culture/2017/11/nfl-coaches-are-making-better-fourth-down-decisions.html>, 2017. Accessed July 20, 2019.
- D. R. Yam and M. J. Lopez. What was lost? a causal estimate of fourth down behavior in the national football league. *Journal of Sports Analytics*, (Preprint):1–15, 2019.