arXiv:2608.27362v1 [stat.AP] 27 Aug 2026

How exceptional was the Big Three era? Extremes and persistence in men's professional tennis

Manuele Leonelli[a]

[a]*School of Science and Technology, IE University, Madrid, Spain*

Abstract

Three players won 66 of the 81 Grand Slam titles contested between 2003 and 2023, and their era is widely held to be the most dominant in the history of tennis. Assessing it means comparing players who never met, so that every comparison passes through the opponents each did face. We measure dominance by how far a player stands above the field of his own day, how many stand there at once, and how long they remain, the last through averages over fixed windows rather than through unbroken runs, which are unstable when the strengths beneath them are themselves estimated. Strengths come from a Bayesian dynamic Bradley-Terry state space model, and dominance is characterized through exceedances of a high threshold. Applied to 197,926 matches played between 1968 and 2025, only Djokovic separates from the field on peak strength, Federer and Nadal being indistinguishable from Borg, McEnroe and Lendl. What distinguishes the recent period is not how high its best players stood but how long three of them stood there together. At its height three players held a top-three position in at least four fifths of all weeks, where in the strongest comparable decade of the past those positions were shared among four. The sharper anomaly lies between: for twelve years from 1990 the upper tail was not merely unoccupied but compressed.

Keywords: OR in sports, extreme value theory, Bradley-Terry models, state space models, uncertainty propagation

1. Introduction

Between Roger Federer's first Grand Slam title at Wimbledon in 2003 and Novak Djokovic's twenty-fourth at the US Open in 2023, Federer, Rafael

Email address: manuele.leonelli@ie.edu (Manuele Leonelli)

Nadal and Djokovic won 66 of the 81 major titles contested, a share of 81.5 per cent, and only ten other players won a major over those two decades (ATP Tour, 2026b). Between them the three occupied the top of the world ranking for 947 weeks, and Djokovic alone accumulated 428 of them, more than any player before him (ATP Tour, 2026a). The period is routinely described as the most dominant by any group in any sport, and comparisons are usually drawn with the rivalry of Björn Borg and John McEnroe in the early 1980s or with the years of Pete Sampras and Andre Agassi in the 1990s. This paper asks what such a comparison requires, and whether the available statistical machinery can deliver it.

A player active in 1980 never meets a player active in 2015, so any statement ranking the two rests on a chain of indirect comparisons mediated by the opponents they did face. Dynamic paired-comparison models construct such chains by converting match outcomes into a latent strength trajectory for each player, and the technology is mature. Our concern is with what happens next. The claim that three players dominated the sport for two decades is not a claim about any one of them. It bundles together three separate questions: how far the best players stood above the rest of the tour, how many of them stood there at the same time, and how long that arrangement lasted. The first asks about the magnitude of an extreme, and the second about how many extremes coincide; for both, the peaks-over-threshold apparatus of extreme value theory supplies the standard tools. The third asks how long a configuration is maintained, and it is the one that usually carries most of the rhetorical weight. It is also the one whose answer depends most sharply on how the question is posed, and part of what follows is concerned with establishing which ways of posing it can be trusted.

The difficulty is easiest to see through an example. Fix a level high enough that only a dominant player clears it, and ask how long three players held it at once. One kind of answer counts weeks: in what fraction of the weeks between 2003 and 2022 were three players simultaneously above the level, and, given that three were above it in some week, how likely is that still to be true a year later? Another measures an unbroken stretch: what was the longest run of consecutive weeks with three players continuously above it? Both are natural descriptions of the same phenomenon and both appear in applied work, but they respond very differently to the smallest disturbance. Suppose that in a single week of 2012 Federer's strength dips just below the level, on a minor injury, an early loss, or nothing more than noise in the rating. The counting answers move by one week in a thousand. The unbroken stretch is halved, a decade of continuous three-way dominance becoming two spells of five years. Averaging over the record absorbs an

isolated interruption; deciding week by week whether a spell has ended does not. We report the first kind throughout and avoid the second. Unbroken stretches are delicate wherever a record must be cut into clusters, and more so when the strengths beneath them are estimated rather than observed; the same difficulty arises for any latent state that is filtered and then thresholded, and sport supplies an unusually complete record on which to examine it.

Answering questions of this kind, about how far above the field, how many at once and for how long, requires a strength for every player in every week, carrying honest uncertainty. We estimate these with a Bayesian dynamic Bradley-Terry state space model (Durbin and Koopman, 2012). Each player is given a latent strength evolving as a random walk. The same model is later given one component per playing surface, correlated across them, so that a result on clay revises what we believe about a player on grass in proportion to how far the surfaces share information. Rather than a single fitted rating, we take samples from the posterior over entire strength trajectories and recompute every quantity within each sample, so that uncertainty about the ratings reaches every number we report. A player counts as dominant in a week if his strength exceeds a threshold, and that threshold is set not in rating points but so that, on average over the whole period, a stated number of players clear it each week. Eras are then compared at a common rate of exceedance rather than a common level. This matters because what is comparable across five decades is not a distance in rating points but how rare it was to stand that far above one's own contemporaries.

Applied to the complete record of men's professional tennis from 1968 to 2025, the picture that emerges is not quite the conventional one. On peak strength only Djokovic separates from the field; Federer and Nadal cannot be distinguished from Borg, McEnroe and Lendl. Nor, on a permissive definition of dominance, were three players dominant at once more often than in the 1980s. What is exceptional is duration, and only for the first half of the era: the probability that three-way dominance survives a full year is 0.44 between 2002 and 2013, against 0.12 in the early 1980s and 0.06 thereafter. The years in between are the sharper anomaly. From 1990 to 2001 the upper tail was not merely unoccupied but compressed, and some of the sense that the era which followed was unprecedented derives from the contrast with the one that preceded it.

2. Literature review

2.1. Paired comparisons and dynamic ratings

The modern treatment of paired-comparison data begins with the model of Bradley and Terry (1952), in which the probability that one player defeats another is a logistic function of the difference between two latent strengths; Cattelan (2012) surveys the subsequent development, with particular attention to the dependence structures that arise when comparisons share players. Extending the static formulation to settings in which strengths evolve is the central methodological problem in competitive ranking, and two broad approaches have emerged. The first is recursive and heuristic, exemplified by the system of Elo (1978) and its descendants, in which ratings are updated after each match by an amount proportional to the discrepancy between the observed and expected result; comparative assessments of such systems in tennis are given by Kovalchik (2016). The second treats strength as a latent stochastic process observed through match outcomes, which places the problem within the framework of nonlinear state space models. Glickman (1999) gave the canonical treatment, approximating the intractable filtering recursions by Gaussian updates and thereby producing ratings accompanied by uncertainty measures. Gorgi et al. (2019) pursue a related high-dimensional dynamic formulation for tennis, and Klaassen and Magnus (2001) provide the point-level evidence that motivates treating match outcomes as informative but noisy realizations of an underlying strength difference.

This literature is well represented in operational research. Baker and McHale (2014) developed a dynamic paired-comparison model with the explicit aim of comparing tennis players across eras. Their solution, a barycentric rational interpolation of time-varying strengths combined with empirical Bayes shrinkage, was applied to the women's tour in Baker and McHale (2017). Further contributions in the same tradition include weighted variants of the Elo recursion (Angelini et al., 2022) and network-based measures of competitive standing (Arcagni et al., 2023). The question we address is therefore not new. We differ from this work in treating the ratings as objects carrying uncertainty that must be propagated, which rules out fitted trajectories such as the barycentric interpolants of Baker and McHale (2014), since a single deterministic trajectory per player carries none, and in asking which downstream summaries survive that propagation.

A parallel literature concerns competitive balance and its evolution. Berry et al. (1999) confronted the era-bridging problem directly, developing a hierarchical model to compare athletes whose careers do not overlap. Radicchi (2011) constructed a prestige measure from the network of matches and used

it to rank players across the professional era, while Breznik et al. (2025) asked whether particular periods constitute eras of concentrated dominance, identifying several through cointegration analysis of ranking trajectories. The latter is the closest antecedent to our substantive question, and we return to the comparison in Section 6.

2.2. Extremes in sport

Characterizing dominance as behaviour in the upper tail connects naturally to extreme value theory, whose general apparatus is set out in Coles (2001). The literature applying these ideas to sport has developed almost entirely within disciplines in which performance is directly measured on a continuous scale. Einmahl and Magnus (2008) estimated ultimate records across athletic events, and the same framework has since been applied to sprinting records in a formulation permitting athlete-specific tails (Einmahl and He, 2026). Stephenson and Tawn (2013) address the related problem of ranking performances through a conceptual population model, and swimming has received careful treatment in Spearing et al. (2021).

A recurring observation in that literature is that dependence structure is typically neglected: analyses either assume that repeated performances by the same athlete are independent, or reduce each athlete to a single best result (Spearing et al., 2023). The framework of Spearing et al. (2023) for extremes of longitudinal data is designed to remedy this and is the closest methodological antecedent to the present work, although it is developed for directly measured performances rather than for latent strengths that must themselves be inferred. That difference matters here, though not only for the reason one might expect. Estimation error does interact with the extremal summaries, but the summaries themselves differ in how much disturbance they can absorb, and the second point is the more fundamental.

The question of whether several extremes occur together has been formalized in the spatial setting by Dombry et al. (2018), who introduced extremal concurrence as the event that the maximum over time is achieved simultaneously at several locations. Related notions appear in the climate literature under the heading of compound and concurrent events (Zscheischler et al., 2020), whose spatially compounding case is the closest to our object, with players playing the role of locations. It differs from the concurrence probability of Dombry et al. (2018) in a respect made explicit in Section 3: we are concerned with sustained joint occupation of the upper tail rather than with the coincidence in time of individual maxima. The two need not agree, since players may hold the tail together for years while peaking at different moments, and we estimate the former rather than the latter.

2.3. Two-step inference

Estimating strengths and then summarizing them is a two-step procedure, and each step has a literature on what can go wrong. The first concerns the summary. A statistical functional is qualitatively robust when small perturbations of the underlying distribution produce small changes in its value (Hampel, 1971), and the longest run of exceedances is not: moving one period across the threshold can halve it. This is why cluster-based summaries are known to be delicate even when the process is fully observed, as the literature on estimating the extremal index, whose reciprocal is the mean cluster size, attests (Ferro and Segers, 2003; Northrop, 2015).

The second concerns the inputs, which here are estimated rather than measured. The difficulty is long recognized for generated regressors (Pagan, 1984) and has become more pressing wherever fitted or predicted values are passed downstream as though they were data (Leonelli and Smith, 2015; Angelopoulos et al., 2023). The remedies proposed in that literature do not transfer here, since the difficulty is not one of variance but of bias and the direction of the bias reverses between estimators. The two difficulties compound: estimation error is exactly the sort of perturbation to which a non-robust summary is sensitive, and a rating scheme supplies it systematically rather than at random. Section 3.4 measures what this costs.

3. Dominance, exceedances and their summaries

This section defines what we mean by dominance and sets out the summaries we shall report. Nothing in it depends on the model of Section 4. We require only some procedure, which may be a recursive rating system, the posterior mean of a statistical model, or a single draw from a posterior, that assigns to each player i and period t a strength $\theta_{i,t}$.

3.1. Relative strength

Absolute strength is not comparable across periods, because the reference level implied by any rating scheme drifts as the population of players turns over. We therefore define the relative strength of player i in period t as

$$r_{i,t} = \theta_{i,t} - \bar{\theta}_t^{(10:100)}, \tag{1}$$

where $\bar{\theta}_t^{(10:100)}$ is the mean strength of the players ranked tenth to hundredth among those active in period t. Excluding the top nine prevents the reference level from being contaminated by the very players whose dominance is under study, while excluding the long tail of marginal participants prevents it from

tracking changes in the size and composition of the professional population. Dominance is thus defined relative to the contemporaneous field rather than on an absolute scale, which is what makes cross-era comparison meaningful.

Relative strengths are reported throughout on the conventional rating scale, obtained by multiplying the latent difference by $400/\log 10$ so that 400 points correspond to ten-to-one odds between two players. A relative strength of 300 therefore means that the player would be expected to win roughly 85% of matches against a player at the level of the reference band. These are not Association of Tennis Professionals (ATP) ranking points, which accumulate results over a season and measure something different. The choice of reference band is a judgement, and Section S4 of the supplementary material reports how the results change when it is widened or narrowed.

3.2. Dominance as a peaks-over-threshold problem

Let u be a threshold and define the exceedance count

$$N_t(u) = \sum_i \mathbf{1}\{r_{i,t} > u\}, \tag{2}$$

the number of players standing more than u above the field in period t. The claim that a sport was dominated by three players amounts to the claim that $N_t(u) \geq 3$ for a high u and for a long time, and everything we report is a summary of the process $\{N_t(u)\}$ or of the exceedances themselves.

This is the peaks-over-threshold construction (Coles, 2001), with two departures from a standard application. The first is that $r_{i,t}$ is latent, which is the subject of the paper. The second is that exceedances are strongly dependent both across players within a period, since $r_{i,t}$ is defined relative to a common field, and across periods within a player, since strength evolves smoothly. It is the second kind of dependence that the summaries below are designed to describe. The magnitudes of the exceedances are reported in Section 5 through order statistics and career peaks rather than through a fitted generalized Pareto tail, since the claims under examination concern how many players clear the threshold and for how long.

One choice has to be made before any of this can be stated precisely. The threshold u requires care, because a latent strength is identified only up to scale, so a threshold fixed in absolute units is not comparable across rating schemes or model specifications. We therefore calibrate u so that the marginal exceedance rate is fixed by construction: given a target mean count $\bar{N}$ over an analysis window of T periods, we set u equal to the $(\bar{N}T)$-th largest value of $r_{i,t}$. When strengths are drawn from a posterior, this calibration is performed separately within each draw. Every comparison we

report is therefore made at a common marginal rate, which removes scale as a confounder and isolates the distributional question of how exceedances are arranged in time. The calibration fixes only the mean of N_t over the window; the probability that three players exceed the threshold at once, and how that event is arranged in time, remain free to differ between eras, and are what is compared. We work throughout with two values of $\bar{N}$, one permissive and one demanding, since the definition of dominance is a matter of degree and the comparison between eras turns out to depend on it in an informative way. Each table states which threshold it uses.

3.3. What we estimate

Three summaries of $\{N_t(u)\}$ are reported throughout. The first is how often the configuration occurs, through the rate $P(N_t \geq c)$ and the expected count $E[N_t]$. The second is the lag persistence probability $\pi_k = P(N_{t+k} \geq c \mid N_t \geq c)$, the chance that a configuration in place now is still in place k periods later. The third is the block persistence probability $\beta_w = P(N_s \geq c \text{ for all } s \in [t, t+w))$, the chance that it holds without interruption throughout a window of w periods. Each is a time average of a bounded function of a fixed-width window, and between them the profiles $\{\pi_k\}$ and $\{\beta_w\}$ convey what a claim about the duration of dominance is usually meant to convey.

The obvious alternative is the longest unbroken run of periods with $N_t \geq c$, and we avoid it deliberately. Summaries of that kind require the record to be partitioned into clusters of exceedance, which is delicate even when the process is observed, as Section 2 noted, because the answer depends on the threshold and on whatever declustering rule separates one cluster from the next. Our setting adds a second difficulty. Strengths here are inferred, so the sequence handed to any such summary carries interruptions of its own, whose direction depends on the estimator: a smooth recursive rating manufactures continuity and lengthens clusters, while independently perturbed posterior samples fragment them. Averages over fixed windows are largely unaffected, since an isolated interruption contributes in proportion to the inverse record length instead of reorganizing the whole answer. Section 3.4 quantifies the difference.

3.4. How the summaries behave

To see what this costs in practice we simulate records for which the answer is known: 250 players over 400 periods, strengths following independent random walks and match outcomes a logistic link, which is the process

Table 1: Behaviour of each functional across 25 simulated records from one process. The first column is the coefficient of variation of the true value across replicates, computed without estimation. The remaining columns give the mean and standard deviation (s.d.) across replicates of the ratio of estimate to truth, so that unity denotes exact recovery, when a recursive rating is treated as observed data and when the posterior median is used. All estimators are calibrated to a common marginal exceedance rate, which fixes the coefficient of variation of the rate functional.

	truth	rating as observed		posterior median	
Functional	c.v.	mean	s.d.	mean	s.d.
Run-based summaries					
Longest run	0.47	1.75	0.87	1.32	0.60
Mean run length	0.60	3.56	2.67	1.79	0.66
Gap-tolerant runs					
Tolerance $g = 4$	0.33	1.35	0.66	1.00	0.35
Tolerance $g = 8$	0.25	1.25	0.47	1.00	0.28
Fixed-window averages					
Lag persistence π_1	0.06	1.12	0.07	1.08	0.07
Lag persistence π_8	0.23	1.41	0.41	1.15	0.29
Rate functional	0.01	1.00	0.01	1.00	0.01

specified in Section 4, with roughly 70 matches per period among 48 active players. We repeat this 25 times and compute every functional three ways on each replicate: from the true strengths, from a recursive Elo rating treated as observed data, and as the posterior median under the scheme of Section 4. All three are calibrated to a common marginal exceedance rate, so no comparison is confounded by the choice of threshold.

Table 1 orders the summaries twice over. The first column does so before any estimator is involved: the longest run and the mean run length vary by roughly half their own size across replicates of one and the same process, so a single history says little about the runs that process tends to produce, whereas π_1 varies by six per cent and the rate functional not at all. The remaining columns show what estimation adds. Treating ratings as observed inflates the two run-based summaries by factors of 1.75 and 3.56, while π_1 and the rate functional are recovered to within a few per cent. Gap-tolerant runs fall between the two, and a tolerance of eight periods removes most of the bias at the price of a coarser description.

Two points are worth adding. Posterior sampling alone does not repair a run-based summary, improving the longest run only from 1.75 to 1.32. The block persistence probabilities, which are levels rather than ratios, are

overstated at every window and most severely at long ones, at 0.132 against a truth of 0.069 for $w = 13$, with the posterior roughly halving the error; this is a reason to read the whole profile rather than any single window.

3.5. Dominance across playing surfaces

Tennis is played on surfaces whose demands differ, so a player may dominate on one and not another, and joint occupation of the upper tail across surfaces is itself a form of exceptional achievement. Writing $r_i^{(s)}$ for a player's career peak relative strength on surface s and $U_i^{(s)}$ for its rank percentile within a reference population, we summarize cross-surface dependence through the coefficients of Ledford and Tawn (1996),

$$\chi(q) = \frac{P(U^{(s)} > q,\, U^{(s')} > q)}{1 - q}, \qquad \eta(q) = \frac{\log(1-q)}{\log P(\min(U^{(s)}, U^{(s')}) > q)}, \tag{3}$$

and we report the number of players simultaneously extreme on all three surfaces, comparing it against the counts implied by independence and by perfect dependence.

4. A dynamic Bradley-Terry state space model

Section 3 requires a strength $\theta_{i,t}$ for every player and period, and this section specifies the model that supplies it.

4.1. Specification

Because tennis is played on several surfaces, we take the state to be a vector rather than a scalar. Let $\boldsymbol{\theta}_{i,t} = (\theta_{i,t}^{(1)}, \ldots, \theta_{i,t}^{(S)})^\top$ denote the latent strengths of player i at time t, one component per surface, where t indexes weekly rating periods. We assume

$$\boldsymbol{\theta}_{i,t} = \boldsymbol{\theta}_{i,t-1} + \mathbf{w}_{i,t}, \qquad \mathbf{w}_{i,t} \sim N_S(\mathbf{0}, \tau^2 \mathbf{R}(\rho)), \tag{4}$$

independently across players, with initial condition $\boldsymbol{\theta}_{i,t_{0i}} \sim N_S(\mathbf{0}, \sigma_0^2 \mathbf{R}(\rho))$ at the period t_{0i} of a player's first appearance. Here $\mathbf{R}(\rho)$ is an exchangeable correlation matrix with off-diagonal ρ, so that a single parameter governs how far improvement on one surface carries over to the others. At $\rho = 1$ the model collapses to a single common strength, and at $\rho = 0$ the surfaces become separate problems. Two specializations are used below. Every result other than those of Section 5.9 takes $S = 1$, a single strength per player estimated from all matches irrespective of surface, which is the appropriate

object when the question concerns dominance over the field as a whole. Section 5.9 takes $S = 3$ and estimates ρ, in order to ask whether occupancy of the upper tail transfers between surfaces.

Conditional on strengths, the probability that player i defeats player j in a match played on surface s over a maximum of f sets is

$$P(i \text{ defeats } j \mid \boldsymbol{\theta}_{i,t}, \boldsymbol{\theta}_{j,t}) = \left[1 + \exp\{-a_f(\theta_{i,t}^{(s)} - \theta_{j,t}^{(s)})\}\right]^{-1}. \tag{5}$$

Men's professional tennis uses two match formats. Matches at the four Grand Slam tournaments are played as the best of five sets, and almost all other matches as the best of three. A longer match affords more opportunity for the stronger player to prevail, so the same strength difference should translate into a larger win probability, and we allow this through the discrimination parameter a_f. We normalize $a_3 = 1$, which fixes the scale of the latent strengths, and estimate a_5, which is identified by the relative frequency of upsets across the two formats. A match thus informs one component of the state directly, and the remaining components indirectly through the correlation in (4).

4.2. The posterior

Write $\boldsymbol{\psi} = (\tau, \sigma_0, \rho, a_5)$ for the hyperparameters, which govern the dynamics and the link and are held fixed at the values estimated in Section 4.3. Given $\boldsymbol{\psi}$ the model has two levels. The strength trajectories $\boldsymbol{\Theta} = \{\boldsymbol{\theta}_{i,t}\}$ follow the Gaussian Markov process (4), which supplies the prior for $\boldsymbol{\Theta}$: it encodes the assumption that strengths change slowly and by small amounts, and it is what allows a player's estimated strength in one week to be informed by his results in neighbouring weeks. The match outcomes $\mathbf{y}$ are then conditionally independent Bernoulli variables with probabilities given by (5). The posterior is

$$p(\boldsymbol{\Theta} \mid \mathbf{y}, \boldsymbol{\psi}) \propto \prod_i \Big[p(\boldsymbol{\theta}_{i,t_{0i}}) \prod_{t > t_{0i}} p(\boldsymbol{\theta}_{i,t} \mid \boldsymbol{\theta}_{i,t-1}) \Big] \prod_m p(y_m \mid \boldsymbol{\theta}_{i_m,t_m}, \boldsymbol{\theta}_{j_m,t_m}), \tag{6}$$

where m indexes matches. Every quantity in Section 3 is a functional of $\boldsymbol{\Theta}$, so all inference proceeds by sampling from (6) and evaluating the functional within each sample.

Two features of (6) prevent this being done directly. The state process is Gaussian but the likelihood is logistic, so no conjugacy is available; and because each match involves two players, the posterior does not factorize across players, so the state is effectively of dimension SN, where N is the

number of players. Exact treatment by Markov chain Monte Carlo remains possible, and we construct such a scheme in Section S3 of the supplementary material, but it does not scale to two hundred thousand matches. We therefore adopt a deterministic approximation, and the same section shows that it reproduces the exact posterior closely on a subsample where both can be run.

4.3. Estimating the hyperparameters

The hyperparameters are treated by empirical Bayes, fixed at the maximizer of the prequential likelihood (Dawid, 1984),

$$\ell(\boldsymbol{\psi}) = \sum_{m} \log p\big(y_m \mid \mathbf{y}_{1:m-1}, \boldsymbol{\psi}\big), \tag{7}$$

restricted to matches in which both players have prior history, since a prediction involving a debutant carries no information about the dynamics. For a state space model this decomposition is the likelihood, subject to the Laplace approximation used in filtering; the predictive probabilities come from the filter of Section 4.4, so maximization is a grid search over $\boldsymbol{\psi}$ with one forward pass at each point. Fixing $\boldsymbol{\psi}$ rather than giving it a prior is itself an approximation, and the component that matters is τ, the weekly innovation standard deviation. We assess it through the curvature of (7) at its maximum and propagate it by refitting at displaced values, with the result reported in Section 5.

4.4. Approximating the posterior

We approximate (6) by a sequence of local Gaussian approximations, following the general strategy of Glickman (1999). Two simplifications are involved. The first is a mean-field step across players: within a rating period each player's strength is updated with the opponents' strengths held at their current values, their uncertainty entering only through an attenuation factor. The second is a Laplace step within each player: the period posterior is replaced by a Gaussian centred at its mode with curvature evaluated there. One point of practice matters. The Glicko rating update takes a single Fisher scoring step towards that mode, which is exact when a player plays one match per period; it is not enough here, since a player may contest seven matches in a fortnight, and in our data the single-step version placed two players of modest achievement among the highest-rated in the sample. We therefore iterate the step to convergence, at negligible cost, and the artefact disappears. The Laplace approximation is then reproduced exactly by a

Gaussian pseudo-observation, which puts the model in linear Gaussian state space form and lets a result on one surface revise the estimated strength on the others in proportion to ρ. The update equations are given in Section S2 of the supplementary material.

4.5. Trajectory sampling

The functionals of Section 3 depend on the joint distribution of the strength process across adjacent periods, not on its pointwise marginals, so pointwise posterior summaries are inadequate. We therefore sample complete trajectories by forward-filtering backward-sampling (Carter and Kohn, 1994). For a player observed in periods $t_1 < \cdots < t_K$ with filtered means $\mathbf{m}_k$ and covariances $\mathbf{P}_k$, we draw $\boldsymbol{\theta}_{t_K} \sim N_S(\mathbf{m}_K, \mathbf{P}_K)$ and then recursively, for $k = K-1, \ldots, 1$,

$$\boldsymbol{\theta}_{t_k} \mid \boldsymbol{\theta}_{t_{k+1}} \sim N_S\big(\mathbf{m}_k + \mathbf{J}_k(\boldsymbol{\theta}_{t_{k+1}} - \mathbf{m}_k),\ \mathbf{P}_k - \mathbf{J}_k\mathbf{P}_k\big)\,,$$

where $\mathbf{J}_k = \mathbf{P}_k(\mathbf{P}_k+\mathbf{Q}_k)^{-1}$ is the smoother gain and $\mathbf{Q}_k = (t_{k+1}-t_k)\tau^2\mathbf{R}(\rho)$ the innovation covariance accumulated across the gap.

5. Eras of dominance on the men's tour

Every quantity below is an estimand of Section 3 computed from the strengths of Section 4, evaluated within each of 300 posterior samples and summarized across them. Results are reported by era, and the division is mechanical: the window is cut into four blocks of twelve years, 1978–1989, 1990–2001, 2002–2013 and 2014–2025. Boundaries chosen to bracket the careers of the players under study would make the comparison less than blind, and blocks of unequal length would give the longer one more opportunity to accumulate whatever is being counted. It happens that 1990 nearly coincides with a genuine break, the replacement of the Grand Prix and World Championship Tennis circuits by the ATP Tour, and that the third block opens the year before Federer's first major title, so the last two between them span the whole period over which the three players overlapped, and the boundary between them splits that period in half so that the two halves can be compared. Because no division of a continuous record is beyond argument, Figure 3 presents the same quantity as a continuous series on which no boundary is imposed, and Section 5.8 compares the strongest decade in the record against the strongest decade disjoint from it, which requires no division at all.

5.1. Data

We analyse the 197,926 men's professional singles matches played on the ATP tour between 1968 and 2025 and recorded in the TML-Database (TennisMyLife, 2026), which takes its results, player identifiers and tournament information from the official ATP records. Each match supplies its date, tournament, surface, format and the identities of the two players; 188,771 of them involve two players who have both appeared before and so contribute to estimation. Matches are aggregated into weekly rating periods. The analysis window runs from 1978 to 2025, because the reference level of (1) rises steeply through the first decade of the open era as the professional population stabilizes and is flat thereafter, lying between 1900 and 1935 on a conventional rating scale in every five-year block from 1978 onwards; it contains 1,924 periods and a median of 558 players active in a period, a player counting as active if he has played within the preceding fifty-two weeks. The surface analysis uses the 143,383 matches played on hard, clay and grass, hard including indoor hard courts, and sets aside carpet, the older indoor surface, which accounted for 17 per cent of matches before 1990 but was withdrawn from the tour and last used in 2017; the rank percentiles of (3) are computed over the 199 players with at least forty matches on each of the three surfaces.

5.2. Model checking

On the 188,771 matches in which both players have prior history the fitted model attains a predictive accuracy of 0.675 and a Brier score of 0.207, in line with published benchmarks for rating-based prediction in tennis (Kovalchik, 2016). The prequential likelihood is sharply peaked in τ, whose posterior standard deviation is under two per cent of its value, and refitting two standard deviations either side leaves every comparison reported below unchanged to two decimal places.

In common with rating models generally, ours is somewhat overconfident, and increasingly so for stronger predictions: in the highest probability band, predictions averaging 0.934 correspond to an observed frequency of 0.910. Recalibration on the logit scale (Platt, 1999), estimated on odd-numbered years and evaluated on even-numbered ones, removes it. The estimated intercept is zero and the slope 0.851, so the model overstates every strength difference by a common factor of about 1.18, and after correction the reliability diagram is flat to within 0.006 in every band. Being a global monotone rescaling, the correction leaves every ranking, every pairwise comparison and every rate-calibrated exceedance set unchanged, and affects only the units in

Table 2: Era-average of the kth largest relative strength, in calibrated rating points above the contemporaneous field, with 95% credible intervals.

Rank	1978–1989	1990–2001	2002–2013	2014–2025
1	397 [376, 424]	264 [246, 284]	364 [347, 387]	377 [350, 400]
2	341 [325, 358]	223 [210, 237]	305 [289, 323]	321 [302, 340]
3	296 [278, 315]	198 [187, 208]	260 [247, 276]	284 [268, 302]
4	256 [239, 272]	179 [170, 188]	229 [214, 246]	237 [221, 257]
5	220 [205, 235]	165 [157, 174]	202 [189, 215]	207 [195, 222]

which strengths are quoted. All reported strengths incorporate it. Selected hyperparameters, the reliability diagram and the details of the calibration are collected in Section S1 of the supplementary material.

5.3. The shape of the top of the field

Before turning to individual players it is useful to see what the top of the field looked like in each era. Table 2 reports, for each rank position from first to fifth, the era-average of the corresponding order statistic of relative strength. This is a time average of a bounded function of the configuration at a single period, and so among the summaries of Section 3 that are reliably estimated.

Two features stand out. The first is that the block of the 1980s is above every later one at every rank. The strongest player of an average week between 1978 and 1989 stood 397 points above the field, against 364 and 377 in the two blocks spanning the careers of Federer, Nadal and Djokovic, and the ordering is the same from first place to fifth. Whatever distinguishes the later period, it is not that the top of the field stood further above the rest. The second is the depth of the gap between those blocks and the one separating them. The strongest player of an average week between 1990 and 2001 stood 264 points above the field, and the fifth-ranked player 165, a level exceeded by the eighth or ninth ranked player in the neighbouring blocks. The whole upper region of the distribution was compressed, not merely its extreme. The credible intervals of the two neighbouring blocks lie entirely above those of 1990–2001 at every rank, so the compression is not an artefact of rating uncertainty.

Figure 1 shows the underlying trajectories, giving for each block the posterior relative strength of the eight players who most often occupied the top three, together with the two rate-calibrated thresholds. It makes the same comparison visible player by player. In the first panel four players cross the stricter threshold but do so in succession rather than together: Borg

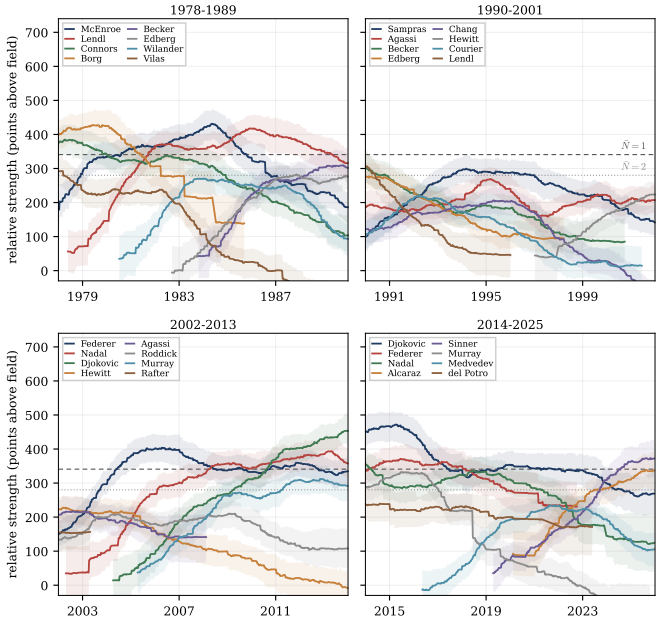


Figure 1: Posterior relative strength of the eight players who most often occupied the top three positions in each block, with 80% credible bands. Horizontal lines mark the two rate-calibrated thresholds, dashed for $\bar{N} = 1$ at 341 points and dotted for $\bar{N} = 2$ at 280 points. All panels share a vertical scale and span twelve years.

is above it from 1978 until 1981, Connors and McEnroe alternate through the early 1980s, and Lendl arrives as the others decline. The second panel contains no trajectory that reaches that line at all; only Sampras spends appreciable time above the looser one, between 1993 and 1997, and the three players who begin the block at about that level, Lendl, Becker and Edberg, fall beneath it within its first two years. The third is the one in which three trajectories sit above the stricter threshold at the same time, from about 2011 until the end of the block. In the fourth Djokovic stands far above it in 2015 and 2016 and then sits at about the line until 2023, while Federer and Nadal fall away beneath it, and the young players who eventually replace them cross it only in the closing years.

The composition of the top three differs more sharply than the levels do. Between 1978 and 1989 six players hold those positions with any regularity: John McEnroe in 74 per cent of weeks, Ivan Lendl 70, Jimmy Connors 56, Björn Borg 36, Boris Becker 23 and Stefan Edberg 21. Between 2002 and 2013 three players account for almost all of it, Federer at 90 per cent, Nadal

Table 3: Posterior peak relative strength over 1978–2025, in calibrated rating points above the contemporaneous reference level. Entries are posterior medians with 95% credible intervals in parentheses, ordered by median reading down the columns.

Player	peak	Player	peak	Player	peak
Djokovic	504 (448, 555)	Connors	411 (366, 468)	Edberg	323 (287, 370)
Borg	461 (412, 523)	Sinner	396 (327, 468)	Wilander	308 (265, 358)
McEnroe	453 (409, 511)	Murray	368 (329, 412)	del Potro	298 (258, 345)
Lendl	449 (401, 507)	Alcaraz	359 (298, 429)	Agassi	294 (254, 336)
Federer	440 (404, 493)	Becker	344 (301, 399)	Hewitt	264 (223, 308)
Nadal	426 (384, 482)	Sampras	333 (291, 381)	Roddick	252 (218, 296)

Table 4: Posterior probability that the row player's peak relative strength exceeds that of the column player.

	Djokovic	Federer	Nadal	Borg	McEnroe	Lendl
Djokovic	–	0.94	0.98	0.85	0.88	0.88
Federer	0.06	–	0.69	0.27	0.36	0.37
Nadal	0.02	0.31	–	0.17	0.25	0.27
Borg	0.15	0.73	0.83	–	0.59	0.64
McEnroe	0.12	0.64	0.75	0.41	–	0.57
Lendl	0.12	0.63	0.73	0.36	0.43	–

74 and Djokovic 51, with Lleyton Hewitt next at 25. Between 2014 and 2025 the concentration is greater still at the top and thinner below: Djokovic holds a top-three position in 98 per cent of weeks, Federer 62 and Nadal 52, with Carlos Alcaraz and Jannik Sinner appearing only towards the end at 29 and 24. The same three men occupy both later blocks, but not in the same way, and Section 5.5 makes the difference precise.

5.4. Peak relative strength

Table 2 describes the top of the field one rank at a time. We turn now to individual players, and in particular to the highest relative strength each attained over his career, which is the quantity the popular debate concerns.

Table 3, Table 4 and Figure 2 report posterior peak relative strength. One player separates from the field. The posterior probability that Djokovic's peak exceeds that of any listed rival lies between 0.85 and 0.98. No other separation is established. In particular, the posterior probability that Federer's peak exceeds Borg's is 0.27, and the corresponding probability for Nadal is 0.17, so that on this criterion the leading players of the early 1980s are, if anything, ahead of two of the three players whose era is usually regarded as

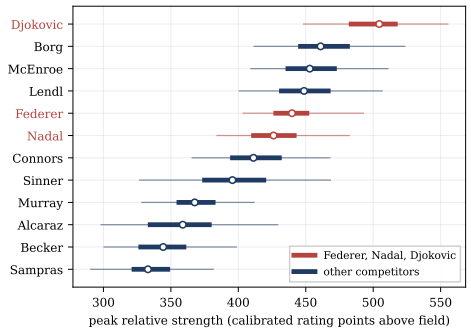


Figure 2: Posterior distribution of peak relative strength. Thick bars denote interquartile ranges and thin bars 95% credible intervals.

Table 5: Concurrent dominance by era at two rate-calibrated thresholds. The threshold is set within each posterior draw so that the mean exceedance count over 1978–2025 equals $\bar{N}$. Total periods with $N_t \geq 3$ are posterior medians with 95% credible intervals.

Era	periods	$E[N_t]$	mean $P(N_t \geq 3)$	periods $N_t \geq 3$	periods $p_t \geq 0.5$
$\bar{N} = 2$ *players per period*					
1978–1989	537	2.79	0.639	341 (250, 434)	315
1990–2001	488	0.42	0.010	0 (0, 28)	0
2002–2013	455	2.21	0.418	191 (138, 233)	194
2014–2025	444	2.56	0.541	241 (163, 330)	217
$\bar{N} = 1$ *player per period*					
1978–1989	537	1.56	0.106	51 (0, 137)	0
1990–2001	488	0.02	0.000	0 (0, 0)	0
2002–2013	455	1.23	0.136	63 (5, 115)	57
2014–2025	444	1.15	0.076	29 (0, 96)	0

exceptional. The credible intervals for Borg, McEnroe, Lendl, Federer and Nadal overlap almost completely.

5.5. Concurrent dominance

Peak strength describes players one at a time. The claim under examination, however, concerns how many of them occupied the upper tail at once, which is the exceedance count $N_t(u)$ of (2). Table 5 and Figure 3 report it by era at the two rate-calibrated thresholds of Section 3. Three findings emerge.

First, at the stricter threshold only one block contains any period in which the simultaneous presence of three dominant players is more probable

Table 6: Persistence profiles by era at the stricter rate-calibrated threshold $\bar{N} = 1$. Entries are posterior medians. The 1990–2001 block admits too few exceedances for the conditional profile to be defined.

	lag persistence π_k					block persistence β_w				
Era	$k=1$	$k=4$	$k=13$	$k=26$	$k=52$	$w=4$	$w=13$	$w=26$	$w=52$	$w=104$
1978–1989	0.89	0.72	0.48	0.28	0.12	0.07	0.03	0.00	0.00	0.00
1990–2001	–	–	–	–	–	0.00	0.00	0.00	0.00	0.00
2002–2013	0.93	0.80	0.67	0.54	0.44	0.11	0.06	0.02	0.00	0.00
2014–2025	0.87	0.66	0.37	0.16	0.06	0.04	0.01	0.00	0.00	0.00

than not, and it is 2002–2013, with 57 such periods against none in any other block, the 1980s included. At the looser threshold three blocks qualify, led by the 1980s with 315 such periods against 217 for 2014–2025 and 194 for 2002–2013.

Second, the two blocks spanning the careers of the three players do not resemble each other. At the looser threshold 2014–2025 is the stronger, 0.541 against 0.418; at the stricter one the ordering reverses decisively, 0.076 against 0.136. The first half had three players far clear of the field, the second had a strong upper order but rarely three players that far clear at once. The 1980s, meanwhile, lead at the looser threshold, 0.639, and trail 2002–2013 at the stricter, 0.106 against 0.136: more players moderately clear of the field, fewer far clear of it. The advantage of 2002–2013 therefore appears only as the definition of dominance is made more demanding, which is the pattern the persistence profiles sharpen below.

Third, the block from 1990 to 2001 is a structural anomaly. Its expected exceedance count is 0.42 at the looser threshold and 0.02 at the stricter one, against values between 1.1 and 2.8 in every other block. On every measure at every threshold in every specification we examined, the upper tail of the relative strength distribution was unoccupied for twelve years.

5.6. Persistence of concurrent dominance

We describe persistence through the lag profile $\{\pi_k\}$ and the block profile $\{\beta_w\}$ rather than through a run length, for the reasons given in Section 3.

Table 6 and Figure 4 report the profiles, and they separate the blocks far more sharply than the rates do. At a lag of one week all three populated blocks agree to within six points, between 0.87 and 0.93. At a lag of one year they do not: the probability that three-way dominance survives is 0.44 for 2002–2013, 0.12 for 1978–1989 and 0.06 for 2014–2025. The block profile tells the same story, with $\beta_{26} = 0.02$ for 2002–2013 against zero elsewhere. The rate functionals of Table 5 do not contain this, their spread across the

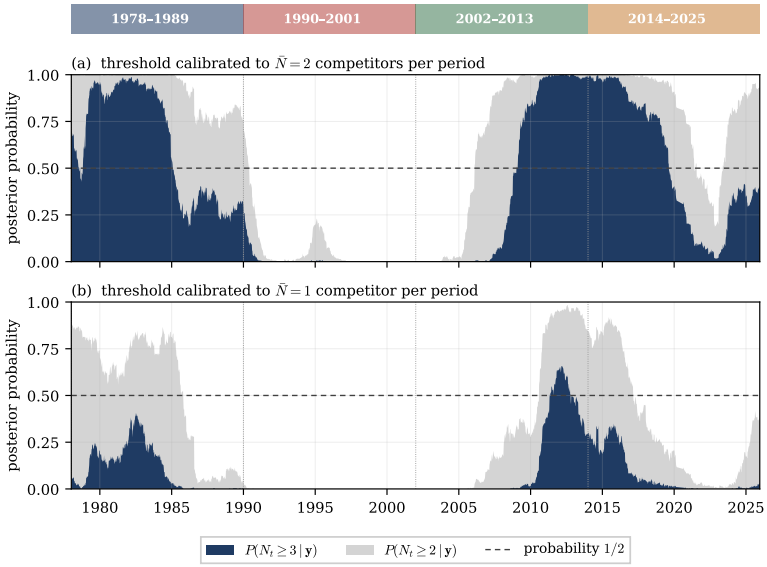


Figure 3: Posterior probability that at least three players (dark) or at least two (grey) are simultaneously dominant, through time, at the two rate-calibrated thresholds. The dashed line marks probability one half; shaded vertical bands denote the four blocks.

same three blocks running only from 0.076 to 0.136: the blocks differ much more in how long dominance was sustained than in how often it occurred, and on that criterion 2002–2013 stands alone, with 2014–2025 below even the 1980s.

5.7. Duration of individual dominance

The profiles above describe the field as a whole. The corresponding quantity for an individual player is his sojourn, the total number of periods in which his relative strength exceeded the threshold, which is a rate functional and therefore among the reliably estimable summaries. Table 7 reports it. The three players of the later era occupy the first three positions; the intervals of the first two exclude the fourth-placed player, while Nadal's overlaps Lendl's. Table 3 orders the same players differently, so the two tables together locate where the later era differs: not in peak strength, but in how long the upper tail was occupied.

Borg records the second-highest peak in the sample and the eighth-longest sojourn, having retired at twenty-six. His interval overlaps Murray's, so the ordering between those two should not be read as settled.

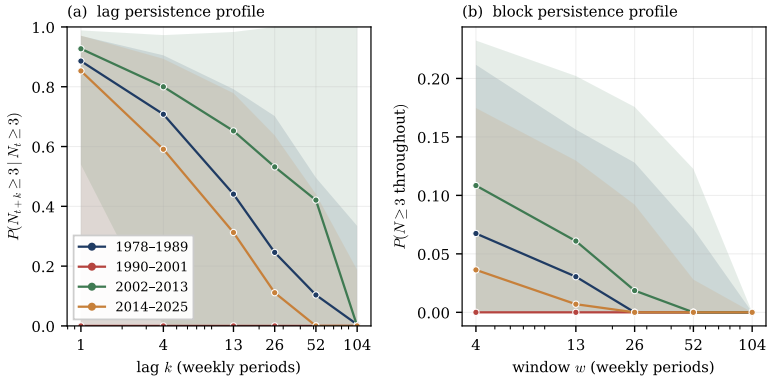


Figure 4: Lag and block persistence profiles by era at the stricter rate-calibrated threshold, with 95% credible bands. Horizontal axes are logarithmic.

Table 7: Posterior sojourn above the dominance threshold at $\bar{N} = 2$, in weekly periods, with 95% credible intervals, ordered reading down the columns.

Player	sojourn	Player	sojourn	Player	sojourn
Federer	612 (526, 714)	Lendl	408 (358, 448)	Murray	215 (118, 302)
Djokovic	578 (457, 657)	McEnroe	365 (298, 439)	Borg	205 (148, 295)
Nadal	510 (396, 598)	Connors	305 (215, 369)	Becker	116 (25, 206)
				Sampras	113 (14, 209)

5.8. The strongest decade against the strongest other decade

The comparison that requires no division into blocks searches over every ten-year window in the record for the one with the highest mean $P(N_t \geq 3)$, then for the best window disjoint from it, and compares the two, so that each period is judged on its strongest decade.

Table 8 reports the outcome, and it is the same at both thresholds. The strongest decade in the record falls in the late 2000s and 2010s, and the strongest decade disjoint from it is the opening of the window, 1978–1987. At the demanding threshold the later decade has twice the mean concurrence, 0.246 against 0.125, and 57 periods in which three-way dominance is more probable than not against none. At the permissive threshold it has 0.911 against 0.708.

The difference is one of kind rather than degree. In 2010–2019 three players held a top-three position in 99, 92 and 80 per cent of weeks, Djokovic, Federer and Nadal, with Murray next at 26. In 1978–1987 four players share the same positions less completely, McEnroe at 84 per cent, Connors 67,

Table 8: The ten-year window with the highest mean $P(N_t \geq 3)$, and the best window disjoint from it, at each rate-calibrated threshold.

	Window	$E[N_t]$	mean $P(N_t \geq 3)$	periods $N_t \geq 3$	periods $p_t \geq 0.5$	π_{52}
$\bar{N} = 1$						
strongest	2010–2019	1.86	0.246	93	57	0.29
best other	1978–1987	1.73	0.125	51	0	0.12
$\bar{N} = 2$						
strongest	2009–2018	3.48	0.911	356	385	0.96
best other	1978–1987	2.93	0.708	324	315	0.80

Lendl 65 and Borg 43. The earlier decade had four strong players trading places; the later had three occupying them at once.

5.9. Dominance across playing surfaces

All results so far pool matches across playing surfaces. We now ask whether occupancy of the upper tail transfers between them, using the coefficients of (3). This requires the surface-specific specialization of (4), with $S = 3$, which we fit afresh and which answers the question in a single fit rather than through three independent per-surface models. Restricting to the 143,383 matches played on the three principal surfaces, involving 4,970 players over 1,818 periods, the cross-surface innovation correlation is estimated at $\rho = 0.90$ at an interior maximum, with $\tau = 0.05$ and predictive accuracy 0.667. The improvement in prequential log-likelihood over the independent specification $\rho = 0$ is 1,365 units, so the correlation is sharply identified. The value implies that about a tenth of the variance of a week's change in strength is specific to the surface on which the results were earned, and since the criterion is predictive, a correlation that did not improve the prediction of matches before they were played would not have been selected, so the approximation of Section 4.4 cannot manufacture it. Strength development is therefore strongly but not completely shared across surfaces, so neither a single common strength nor three separate analyses would be adequate. We summarize career peak strengths through their ranks within the population of 199 players with at least forty matches on each surface, and compute the dependence coefficients of Ledford and Tawn (1996) within each posterior draw.

Tables 9 and 10 and Figure 5 report the results. Upper tail dependence is substantial but incomplete, and the weakest link is between clay and grass, where χ falls to 0.40 at the more demanding level. The number of players extreme on all three surfaces is eight at $q = 0.90$ and four at $q = 0.95$, against expectations of 0.20 and 0.02 under independence and 20 and 10

Table 9: Cross-surface upper tail dependence under the joint model, posterior medians with 95% credible intervals.

	$q = 0.90$		$q = 0.95$	
Pair	χ	η	χ	η
Hard–clay	0.60 (0.50, 0.70)	0.82 (0.77, 0.87)	0.60 (0.40, 0.70)	0.86 (0.77, 0.89)
Hard–grass	0.65 (0.50, 0.75)	0.84 (0.77, 0.89)	0.60 (0.40, 0.80)	0.86 (0.77, 0.93)
Clay–grass	0.50 (0.35, 0.60)	0.77 (0.69, 0.82)	0.40 (0.30, 0.60)	0.77 (0.71, 0.86)
Extreme on all three	8 (6, 11)		4 (3, 5)	
Under independence	0.20		0.02	
Under perfect dependence	20		10	

Table 10: Posterior probability of being extreme on all three surfaces simultaneously, and posterior median surface rank percentiles, under the joint model.

	P(extreme on all three)		median rank percentile		
Player	$q = 0.90$	$q = 0.95$	hard	clay	grass
Djokovic	1.00	1.00	0.995	0.990	0.995
Federer	1.00	1.00	0.990	0.985	0.990
Nadal	1.00	0.97	0.978	0.995	0.980
Murray	1.00	0.91	0.980	0.970	0.985
McEnroe	0.70	0.12	0.960	0.925	0.968
del Potro	0.67	0.12	0.955	0.950	0.930
Lendl	0.57	0.12	0.980	0.975	0.915
Hewitt	0.32	0.01	0.930	0.885	0.960
Soderling	0.24	0.01	0.905	0.940	0.895
Sampras	0.23	0.01	0.970	0.870	0.965

under perfect dependence. At the more demanding level, four players have posterior probability above one half of occupying the upper tail on every surface: Djokovic and Federer at unity, Nadal at 0.97 and Murray at 0.91. Surface universality is therefore a rare property, shared by all three of the players under study together with one contemporary, and by no predecessor with probability above 0.12.

Fitting three separate models to the same data understates this conclusion. Under separate fits Nadal's posterior probability at $q = 0.95$ was 0.32 rather than 0.97, and his grass percentile 0.940 rather than 0.980, because grass matches are the scarcest of the three and an independent fit cannot borrow the information his hard and clay records supply. With ρ estimated at 0.90 the borrowing is substantial and the estimates are correspondingly sharper.

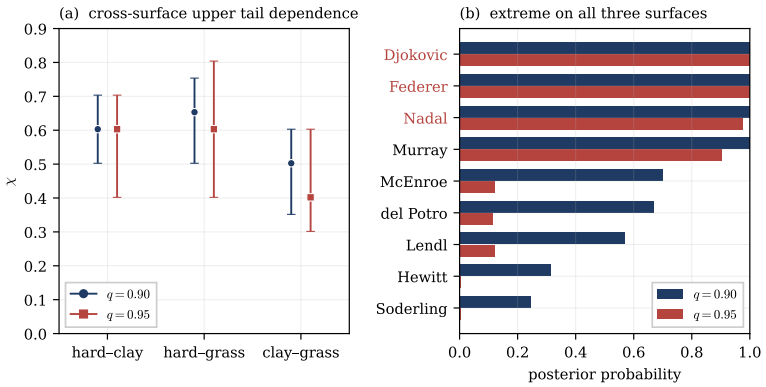


Figure 5: Left: posterior cross-surface upper tail dependence coefficients under the joint model, with credible intervals. Right: posterior probability of occupying the upper tail on all three surfaces simultaneously.

6. Discussion

6.1. Comparison with previous findings

Three earlier answers to the question of the strongest player of the open era can be set against Table 3. Baker and McHale (2014) found Federer the most likely candidate, with Borg and Connors close behind, and noted that their rankings largely agreed with the count of major titles. Their record closes in the early 2010s, before most of Djokovic's career, so the one separation we establish could not have appeared in it; Borg, on our evidence, is if anything ahead of Federer rather than behind him, with posterior probability 0.73. The agreement with title counts is where we part company. Borg, McEnroe and Lendl won 26 majors between them against 66 for Federer, Nadal and Djokovic, yet on peak the five players other than Djokovic cannot be told apart. A title count accumulates over a career, so it is closer to a sojourn than to a peak, and on sojourn our ordering does agree with it. The same distinction accounts for Radicchi (2011), whose cumulative prestige score placed Connors first, Federer seventh and Borg tenth: Borg is second on peak and eighth on sojourn, having retired at twenty-six, and Federer fifth and first. The lists disagree because the quantities do. Baker and McHale (2017) found for the women's tour that intervals on the rankings left several players in contention, and Table 4 says the same of the men.

The closest antecedent on eras is Breznik et al. (2025), who identify five strong eras over a little more than thirty years of data, the most recent ending in July 2019, with Federer the player present in most of them. On

both points we agree: the strongest decade in Table 8 ends in 2018 or 2019 depending on the threshold, and Federer's sojourn is the longest in Table 7. Their record does not reach the early 1980s, however, which is where our evidence places the precedent, at least in how many players were dominant at once. An era in their analysis is, moreover, either present or absent, whereas the distinction that matters most here, between 2002–2013 and 2014–2025 in the chance that three-way dominance survives a year, is one of degree.

Bayram and Barla (2026) report from the same match record that the sport has become more competitive, with more frequent upsets and top-player dominance spread more evenly across the field. This is not in tension with our findings once threshold dependence is taken into account: the body of the distribution can compress while its extreme tail concentrates, and at the permissive threshold the 1980s lead every later block while at the demanding one 2002–2013 does. Table 2 adds that the change was not monotone, since the top of the field stood lower above the rest in 1990–2001 than in any block before or since. Radicchi (2011) named Sampras the best player of 1991–2000, and his peak of 333 points, more than a hundred below the leaders of the blocks on either side, is that compression seen from the top.

Finally, our comparisons are of gaps, not levels. Berry et al. (1999) bridge eras in order to compare all players directly, which requires assumptions about how the talent pool changes; the relative strength of (1) measures each player against the tenth to hundredth ranked players of his own day. That the top of the field stood higher above the rest in the 1980s than in the 2000s is therefore compatible with the later field being stronger throughout, as the increase in competitiveness just noted suggests. Whether Borg at his peak would have beaten Djokovic at his is not a question the match record settles, and we have not asked it.

6.2. Extremes of inferred strengths

Extreme value methods in sport have been applied to measured performances: ultimate records (Einmahl and Magnus, 2008; Einmahl and He, 2026), rankings of performances (Stephenson and Tawn, 2013) and swim times (Spearing et al., 2021). Tennis produces no such measurement. A match yields a winner, and relative strength is a latent variable of a statistical model; we are not aware of a previous application of the peaks-over-threshold apparatus to a sport of that kind. Three things change: the threshold is calibrated to a rate, so that eras are compared at equal rarity; exceedances are dependent across players by construction, and the count $N_t(u)$ absorbs that dependence into a single process, the estimand being sustained joint

occupation of the tail rather than the concurrence of maxima, as Section 2 noted; and the dependence within a player that Spearing et al. (2023) model directly for measured performances is here induced by the state equation (4), so that its standard tail summary, the extremal index, cannot be estimated from inferred strengths.

The last point is easily misread, since estimation aggravates the difficulty without causing it. The longest run is not qualitatively robust (Hampel, 1971): it varies by half its own size across replicates of one process even when strengths are known exactly, and estimation then supplies the interruptions to which it is sensitive. Posterior sampling reduces the inflation only from 1.75 to 1.32, so the remedy is a change of summary rather than of estimator. Once it is made, Section S5 of the supplementary material shows that on the tennis data the value obtained by treating ratings as observed falls inside the posterior interval at both thresholds.

6.3. Limitations

Four limitations deserve statement. First, sojourn above a dominance threshold conflates playing strength with career duration, and career duration depends on factors, including advances in the medical management of athletes, that lie outside the model. The comparison between Borg's peak and his sojourn is a direct illustration; the limitation is one of the estimand rather than of estimation, and it should temper the reading of Table 7. Second, our analysis concerns the men's professional tour. The motivating question is specific to that population, and we make no claim about whether the same structure obtains elsewhere. Third, the reference level in (1) is one of several defensible choices. Section S4 of the supplementary material reports the effect of widening and narrowing the excluded group: peak strengths shift by up to 20 points, but no ordering and no era comparison is disturbed. Fourth, the cross-surface correlation is held constant over a period in which the surfaces themselves have converged in pace (Bayram and Barla, 2026); a correlation allowed to drift would probably rise over the record, and the results of Section 5.9 should be read with that in mind.

7. Conclusions

Applied to the complete record of men's professional tennis, the analysis supports the conventional account of the recent era in one respect and not in others. Only Djokovic separates from the field on peak strength, Federer and Nadal being indistinguishable from Borg, McEnroe and Lendl; what distinguishes the period is how long three players occupied the upper tail together,

and even that holds for 2002–2013 rather than throughout, the probability that three-way dominance survives a year being 0.44 then against 0.12 in the early 1980s and 0.06 between 2014 and 2025. The sharper anomaly is the twelve years from 1990, in which the upper tail was not merely unoccupied but compressed. Each rests on summaries that average over fixed windows rather than on the run lengths that usually carry such claims, and Section 3.4 is our reason for insisting on the difference. The structure is not peculiar to sport. A volatility filtered from asset returns or the condition of a machine inferred from sensor readings is a latent process that is thresholded and then described by how long it stayed beyond the threshold, and the same choice between run-based and fixed-window summaries arises in each. Whether the distinction holds as generally as that is the question we would take up next.

Data and code

All results are computed from publicly available match records. Data and code to replicate every table and figure are available at `https://github.com/manueleleonelli/Tennis_Extremes`.

Acknowledgements

This work was supported by the Spanish Agencia Estatal de Investigación grant PID2023-153222OB-I00.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work the author used generative AI tools to support coding and writing. After using these tools, the author reviewed and edited the content as needed and takes full responsibility for the content of the published article.

References

Angelini, G., Candila, V., De Angelis, L., 2022. Weighted Elo rating for tennis match predictions. European Journal of Operational Research 297, 120–132.

Angelopoulos, A.N., Bates, S., Fannjiang, C., Jordan, M.I., Zrnic, T., 2023. Prediction-powered inference. Science 382, 669–674.

Arcagni, A., Candila, V., Grassi, R., 2023. A new model for predicting the winner in tennis based on the eigenvector centrality. Annals of Operations Research 325, 615–632.

ATP Tour, 2026a. Federer, Nadal and Djokovic: the impact of the legendary Big Three. https://www.atptour.com/en/news/atp-no-1-club-docuseries-part-4-big-3-feature.

ATP Tour, 2026b. Grand Slams: tournaments, records and statistics. https://www.atptour.com/en/news/grand-slams-tournaments-records-stats.

Baker, R.D., McHale, I.G., 2014. A dynamic paired comparisons model: Who is the greatest tennis player? European Journal of Operational Research 236, 677–684.

Baker, R.D., McHale, I.G., 2017. An empirical Bayes model for time-varying paired comparisons ratings: Who is the greatest women's tennis player? European Journal of Operational Research 258, 328–333.

Bayram, F., Barla, A., 2026. A longitudinal analysis of generational shifts in professional men's tennis across the open era. Scientific Reports .

Berry, S.M., Reese, C.S., Larkey, P.D., 1999. Bridging different eras in sports. Journal of the American Statistical Association 94, 661–676.

Bradley, R.A., Terry, M.E., 1952. Rank analysis of incomplete block designs: I. The method of paired comparisons. Biometrika 39, 324–345.

Breznik, K., Candila, V., Milekhina, A., Restaino, M., 2025. Eras of dominance: identifying strong and weak periods in professional tennis. Computational Statistics 40, 2049–2066.

Carter, C.K., Kohn, R., 1994. On Gibbs sampling for state space models. Biometrika 81, 541–553.

Cattelan, M., 2012. Models for paired comparison data: a review with emphasis on dependent data. Statistical Science 27, 412–433.

Coles, S., 2001. An Introduction to Statistical Modeling of Extreme Values. Springer, London.

Dawid, A.P., 1984. Present position and potential developments: some personal views. Statistical theory: the prequential approach. Journal of the Royal Statistical Society: Series A 147, 278–292.

Dombry, C., Ribatet, M., Stoev, S., 2018. Probabilities of concurrent extremes. Journal of the American Statistical Association 113, 1565–1582.

Durbin, J., Koopman, S.J., 2012. Time Series Analysis by State Space Methods. 2 ed., Oxford University Press, Oxford.

Einmahl, J.H.J., He, Y., 2026. Accurate estimates of ultimate 100-meter records. Extremes .

Einmahl, J.H.J., Magnus, J.R., 2008. Records in athletics through extreme-value theory. Journal of the American Statistical Association 103, 1382–1391.

Elo, A.E., 1978. The Rating of Chessplayers, Past and Present. Arco, New York.

Ferro, C.A.T., Segers, J., 2003. Inference for clusters of extreme values. Journal of the Royal Statistical Society: Series B 65, 545–556.

Glickman, M.E., 1999. Parameter estimation in large dynamic paired comparison experiments. Journal of the Royal Statistical Society: Series C 48, 377–394.

Gorgi, P., Koopman, S.J., Lit, R., 2019. The analysis and forecasting of tennis matches by using a high dimensional dynamic model. Journal of the Royal Statistical Society: Series A 182, 1393–1409.

Hampel, F.R., 1971. A general qualitative definition of robustness. The Annals of Mathematical Statistics 42, 1887–1896.

Klaassen, F.J.G.M., Magnus, J.R., 2001. Are points in tennis independent and identically distributed? Evidence from a dynamic binary panel data model. Journal of the American Statistical Association 96, 500–509.

Kovalchik, S.A., 2016. Searching for the GOAT of tennis win prediction. Journal of Quantitative Analysis in Sports 12, 127–138.

Ledford, A.W., Tawn, J.A., 1996. Statistics for near independence in multivariate extreme values. Biometrika 83, 169–187.

Leonelli, M., Smith, J.Q., 2015. Bayesian decision support for complex systems with many distributed experts. Annals of Operations Research 235, 517–542.

Northrop, P.J., 2015. An efficient semiparametric maxima estimator of the extremal index. Extremes 18, 585–603.

Pagan, A., 1984. Econometric issues in the analysis of regressions with generated regressors. International Economic Review 25, 221–247.

Platt, J.C., 1999. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods, in: Advances in Large Margin Classifiers. MIT Press, pp. 61–74.

Polson, N.G., Scott, J.G., Windle, J., 2013. Bayesian inference for logistic models using Pólya-Gamma latent variables. Journal of the American Statistical Association 108, 1339–1349.

Radicchi, F., 2011. Who is the best player ever? A complex network analysis of the history of professional tennis. PLoS ONE 6, e17249.

Spearing, J., Tawn, J., Irons, D., Paulden, T., 2023. A framework for statistical modelling of the extremes of longitudinal data, applied to elite swimming. arXiv preprint arXiv:2306.12419 .

Spearing, J., Tawn, J., Irons, D., Paulden, T., Bennett, G., 2021. Ranking, and other properties, of elite swimmers using extreme value theory. Journal of the Royal Statistical Society: Series A 184, 368–395.

Stephenson, A.G., Tawn, J.A., 2013. Determining the best track performances of all time using a conceptual population model for athletics records. Journal of Quantitative Analysis in Sports 9, 67–76.

TennisMyLife, 2026. TML-Database: complete live ATP match database. `https://stats.tennismylife.org/tennis-match-database`.

Zscheischler, J., Martius, O., Westra, S., Bevacqua, E., Raymond, C., et al., 2020. A typology of compound weather and climate events. Nature Reviews Earth and Environment 1, 333–347.

Appendix A. Fitted hyperparameters and calibration

Table A.11: Selected hyperparameters of the model of Section 4 of the main paper.

Innovation standard deviation τ	0.0473
Posterior standard deviation of τ	0.00080
Initial standard deviation σ_0	1.0
Discrimination for best-of-five matches a_5	1.15
Cross-surface innovation correlation ρ	0.90
Recalibration slope	0.851
Posterior draws	300

Table A.11 collects the values selected by the procedure of Section 4.3 of the main paper. Fitting a quadratic to a fine grid in τ gives a mode of 0.0473 with an implied posterior standard deviation of 0.00080 and a 95% interval of $[0.0457, 0.0489]$, so that any weakly informative prior would be numerically irrelevant at this sample size. The discrimination for best-of-five matches is estimated at 1.15, confirming that the longer match is more discriminating, though the improvement in fit is modest.

The overconfidence noted in Section 5.2 of the main paper might be addressed by adding a dispersion parameter to the link, equation (5) there, and it is worth recording why this does not work. Such a parameter is identified only jointly with the scale of the latent process, so the optimization runs along a ridge, gaining 0.002 log-likelihood units per match while migrating without settling. Recalibration after fitting diagnoses the same quantity cleanly. Figure A.6 shows the reliability diagram before and after.

Appendix B. The within-period update

Section 4.4 of the main paper approximates the posterior by a mean-field step across players and a Laplace step within each player. This section gives the resulting update equations. Because each match involves a single component of the state, the within-period problem is univariate. Write $m_{i,t}$ and $P_{i,t}$ for the prior mean and variance of the relevant component, carried into period t after the propagation implied by equation (4) of the main paper. Let y_{ij} equal one if i beat j and zero otherwise, and let

$$g_j = a_f\big(1 + 3P_{j,t}/\pi^2\big)^{-1/2}$$

be the attenuation induced by uncertainty about the opponent's strength, the device introduced by Glickman (1999) for propagating an approximately

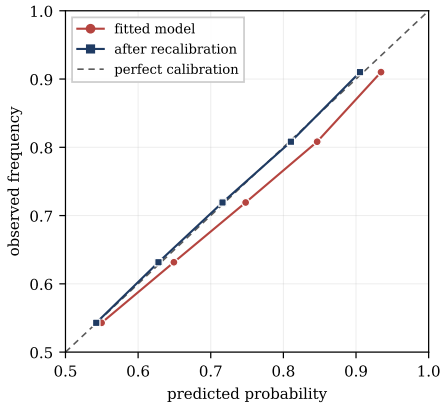


Figure A.6: Reliability of the fitted model before and after recalibration on the logit scale. The estimated correction is a pure rescaling, with intercept zero and slope 0.851.

Gaussian belief through a logistic link. The mode $\hat{\theta}_{i,t}$ of the period posterior solves

$$S_{i,t}(\theta) = \frac{\theta - m_{i,t}}{P_{i,t}}, \tag{B.1}$$

where $S_{i,t}(\theta) = \sum_j g_j\{y_{ij} - E_{ij}(\theta)\}$ is the score contributed by the matches i played in period t, and $E_{ij}(\theta) = [1 + \exp\{-g_j(\theta - \hat{\theta}_{j,t})\}]^{-1}$ is the implied win probability against opponent j.

Equation (B.1) has no closed-form solution. The Glicko rating update takes a single Fisher scoring step from the prior mean, whose accuracy degrades when $S_{i,t}$ is large; as noted in Section 4.4 of the main paper, we solve (B.1) by Newton's method instead. Writing $I_{i,t} = \sum_j g_j^2 E_{ij}(1 - E_{ij})$ for the observed information at the converged mode, the period posterior variance is $\hat{P}_{i,t} = (P_{i,t}^{-1} + I_{i,t})^{-1}$.

To recover a linear Gaussian representation suitable for smoothing, note that the Laplace approximation to the period posterior is reproduced exactly by a Gaussian pseudo-observation $y_{i,t}$ with variance $v_{i,t} = I_{i,t}^{-1}$, where

$$y_{i,t} = \hat{\theta}_{i,t} + \frac{v_{i,t}}{P_{i,t}}\big(\hat{\theta}_{i,t} - m_{i,t}\big).$$

Writing $\mathbf{K} = \mathbf{P}_{i,t}\mathbf{e}_s\,(\mathbf{e}_s^\top \mathbf{P}_{i,t}\mathbf{e}_s + v_{i,t})^{-1}$ for the Kalman gain, where $\mathbf{e}_s$ selects the surface on which the match was played, the full state is updated by

$$\mathbf{m}_{i,t} \leftarrow \mathbf{m}_{i,t} + \mathbf{K}\big(y_{i,t} - \mathbf{e}_s^\top \mathbf{m}_{i,t}\big), \qquad \mathbf{P}_{i,t} \leftarrow \mathbf{P}_{i,t} - \mathbf{K}\,\mathbf{e}_s^\top \mathbf{P}_{i,t},$$

so that a result on one surface revises the estimated strength on the others in proportion to ρ. Under the scalar specialization this step is vacuous and the recursion reduces to the univariate update above.

Appendix C. Accuracy of the posterior approximation

The approximations of Section 4 of the main paper replace an intractable posterior with a tractable one, and their adequacy is an empirical question. We answer it by constructing an exact sampler for the posterior, equation (6) of the main paper, and comparing the two on a subsample small enough for the exact scheme to be run to convergence.

The exact scheme uses the Pólya-Gamma data augmentation of Polson et al. (2013), written PG below. Introducing latent variables $\omega_m \sim \mathrm{PG}(1, \psi_m)$ with $\psi_m = \theta_{i_m,t_m} - \theta_{j_m,t_m}$ renders the logistic likelihood conditionally Gaussian in the strengths, since

$$\frac{\exp(\kappa_m \psi_m)}{1+\exp(\psi_m)} \;\propto\; E_{\omega_m} \exp\left\{\kappa_m \psi_m - \tfrac{1}{2}\omega_m \psi_m^2\right\}, \qquad \kappa_m = y_m - \tfrac{1}{2}.$$

Conditionally on $\{\omega_m\}$ and on the trajectories of all other players, each player's trajectory follows a linear Gaussian state space model with pseudo-observations $\theta_{j_m,t_m} + \kappa_m/\omega_m$ of precision ω_m, and can be drawn exactly by the same backward-sampling recursion used above. Cycling over players in random order and alternating with the Pólya-Gamma draw yields a Gibbs sampler targeting (6) without approximation. The scheme makes no mean-field simplification: the coupling across players induced by (6) is respected exactly, at the cost of requiring many sweeps.

We apply this to a subsample of the tennis records, described in full in Section 5.1 of the main paper, comprising all matches between 2000 and 2010 among players with at least 120 matches in that window. This gives 21,163 matches among 177 players over 422 weekly periods. We run 1,200 iterations and discard the first 400.

Table C.12 reports the comparison. Agreement is close. Posterior means correlate at 0.997 across player-periods and differ by 7.2 rating points on average, against posterior standard deviations of roughly 45 points, and the approximate posterior reproduces the exact posterior spread to three decimal places.

One qualification applies. On the raw scale of the latent strengths the approximate posterior means are displaced downward relative to the exact ones by roughly 40 points, an artefact of the mean-field treatment of opponents, and the approximate posterior standard deviations are about six per

Table C.12: The approximate posterior compared against the exact Pólya-Gamma sampler. Quantities are computed on the relative scale of equation (1) of the main paper, which is the scale on which all results are reported.

Correlation of posterior means across player-periods	0.997
Mean absolute difference in posterior mean	7.2 points
Ninetieth percentile of absolute difference	15.2 points
Ratio of posterior standard deviations, approximate to exact	1.000
Correlation of posterior medians of peak strength	0.998
Mean absolute difference in posterior median of peak strength	6.6 points
Ratio of posterior standard deviations of peak strength	1.007

cent too small. Both effects are almost entirely common across players and therefore cancel under the centring in (1). This is a further reason, beyond the drift in the reference level discussed in Section 5.1 of the main paper, to conduct the analysis on relative rather than absolute strengths.

Appendix D. Sensitivity analyses

The reference level in equation (1) of the main paper excludes the nine strongest active players so that the very players under study do not set the standard against which they are judged. The choice is a compromise: excluding too few contaminates the reference, while excluding too many measures dominance against a field that excludes everyone who might plausibly be called dominant. We therefore repeat the analysis with the reference band running from the fifth to the hundredth ranked player, which excludes only four, and from the fifteenth to the hundredth, which excludes fourteen.

Table D.13 reports the comparison. Widening the excluded group raises every peak relative strength, as it must, since a reference computed from weaker players sits lower. The shift is close to uniform, between 7 and 11 points per step, so the orderings are unaffected: the posterior probability that Djokovic's peak exceeds Federer's moves from 0.92 to 0.94, and the probability that Federer's exceeds Borg's from 0.28 to 0.26.

The concurrence and persistence results are equally stable, and at the looser threshold essentially invariant: the mean probability of three-way dominance moves from 0.639 to 0.640 in 1978–1989, from 0.417 to 0.418 in 2002–2013 and from 0.545 to 0.538 in 2014–2025. At the stricter threshold 1978–1989 gives 0.105 to 0.106, 2002–2013 0.128 to 0.139 and 2014–2025 0.072 to 0.082. The one-year lag persistence moves from 0.41 to 0.47 in 2002–2013 and from 0.11 to 0.13 in 1978–1989, so the ratio between them,

which carries the substantive claim, lies between 3.6 and 3.7 under every band. Excluding fewer players narrows the gap slightly, as expected, since the fifth to ninth ranked players of a strong era are themselves comparatively strong and raise the reference. The effect is small enough not to disturb any conclusion.

A second choice worth testing is the format effect of equation (5) of the main paper. Because best-of-five matches are more discriminating, a given win rate implies a smaller strength difference when achieved over five sets than over three, so an era containing proportionally more of them would have its strengths compressed. The mix does drift, and it drifts against the later era rather than in its favour. Refitting with $a_5 = 1$, which removes the effect entirely, raises every relative strength by roughly three per cent, but does so almost uniformly across periods: averaged over the top five positions the increase is close to 3.2 per cent in every block, the spread across blocks being about a tenth of a percentage point. The differential between eras is thus around a tenth of a percentage point, against gaps of several per cent between the quantities being compared. The format effect therefore acts as a global rescaling rather than as an advantage to any period. This is not because the mix is constant, since it is not, but because the effect is small: fifteen per cent additional discrimination applied to roughly a fifth of matches leaves relative strengths moving by well under one per cent even when the mix shifts by several percentage points. The likelihood prefers $a_5 = 1.15$, but only slightly, the predictive accuracy falling from 0.6745 to 0.6742 when the effect is removed.

Appendix E. Comparison against treating ratings as observed

Section 3.4 of the main paper established how the functionals behave when the truth is known. It remains to ask what difference propagation makes on the tennis data themselves. Table E.14 compares the two. Once the marginal rate is held fixed, the value obtained by treating ratings as observed falls inside the posterior interval for both functionals at both thresholds, and at the looser threshold the two agree within four per cent on the rate functional. On these data, therefore, treating ratings as observed is adequate.

This could not have been established in advance. Section 3.4 of the main paper shows that the longest run is liable to twofold inflation, and nothing in the data indicates which regime obtains. Propagation did not overturn a conclusion here; it established that one could be drawn.

Table D.13: Sensitivity to the reference band used in equation (1) of the main paper. Peak relative strengths are posterior medians in calibrated rating points; the remaining rows are posterior medians of the quantities named.

Reference band	5–100	10–100	15–100
Peak relative strength			
Djokovic	493	504	513
Borg	451	461	468
McEnroe	444	453	460
Lendl	439	449	456
Federer	432	440	447
Nadal	415	426	434
Sampras	325	333	340
Posterior probability of ordering			
P(Djokovic > Federer)	0.92	0.94	0.94
P(Djokovic > Borg)	0.83	0.85	0.84
P(Federer > Borg)	0.28	0.27	0.26
P(Nadal > Borg)	0.17	0.17	0.17
Mean $P(N_t \geq 3)$*, threshold* $\bar{N} = 1$			
1978–1989	0.105	0.106	0.106
1990–2001	0.000	0.000	0.000
2002–2013	0.128	0.136	0.139
2014–2025	0.072	0.076	0.082
Mean $P(N_t \geq 3)$*, threshold* $\bar{N} = 2$			
1978–1989	0.639	0.639	0.640
1990–2001	0.011	0.010	0.009
2002–2013	0.417	0.418	0.418
2014–2025	0.545	0.541	0.538
Lag persistence π_{52}*, threshold* $\bar{N} = 1$			
1978–1989	0.11	0.12	0.13
1990–2001	–	–	–
2002–2013	0.41	0.44	0.47
2014–2025	0.03	0.06	0.07

Table E.14: Ratings treated as observed, against full posterior propagation on the tennis data, both calibrated to a common marginal exceedance rate.

	$\bar{N} = 1$		$\bar{N} = 2$	
Functional	as observed	posterior	as observed	posterior
Longest run of $N_t \geq 3$	39	52 (19, 112)	359	293 (178, 420)
Total periods $N_t \geq 3$	127	156 (80, 249)	815	784 (683, 872)