arXiv:2611.10152v1 [stat.AP] 6 Nov 2026

Modeling the Joint Distribution of Box Score Statistics: A Gaussian Copula Approach

Kenny Watts
University of Virginia
xen5ek@virginia.edu

November 2026

Abstract

What is the probability that LeBron James records at least 27 points, 7 rebounds, and 7 assists tonight? While an interesting statistical question by itself, the rise of sports betting and multi-leg parlays has made answering this question increasingly relevant. However, coming up with an estimated likelihood is tricky, since box score statistics aren't independent.

We estimate the joint probability that a player reaches a combination of thresholds in points, assists, and rebounds using a Gaussian copula, fit separately for each player using only their previous games. We compare it with two benchmarks: the product of the empirical marginal probabilities (i.e. assuming independence) and the empirical frequency of the joint event. Across $\sim 22,000$ player-observations, the Gaussian copula performs on par or better than either benchmark at lower threshold levels, with large gains at high joint exceedance levels. Because it is generative, the model also allows for probabilistic calculations for combinations that a player has rarely or never produced.

1 Introduction

Same-game parlays allow bettors to combine several outcomes from a single game into one wager. In basketball these legs often involve player box score statistics; for example, a bettor may wager that a player records at least 20 points, 5 rebounds, and 5 assists. There is evidence that these multi-leg parlays are far more profitable for sports books (Whelan, 2026), likely due to the difficulty in estimating the probability of the joint event and the multiplicative way that sports books set their parlay odds (every line has a built in vig, so combining multiple bets multiplies the house's take).

Treating the three box score statistics as independent is a poor approximation: a player who scored a large number of points likely spent more time on the court, accumulating opportunities to record rebounds and assists. Modeling these joint probabilities therefore requires a framework that accounts for inter-variable dependencies, rather than treating each metric in isolation. Figure 1 shows that our three box score statistics exhibit mild but positive correlation.

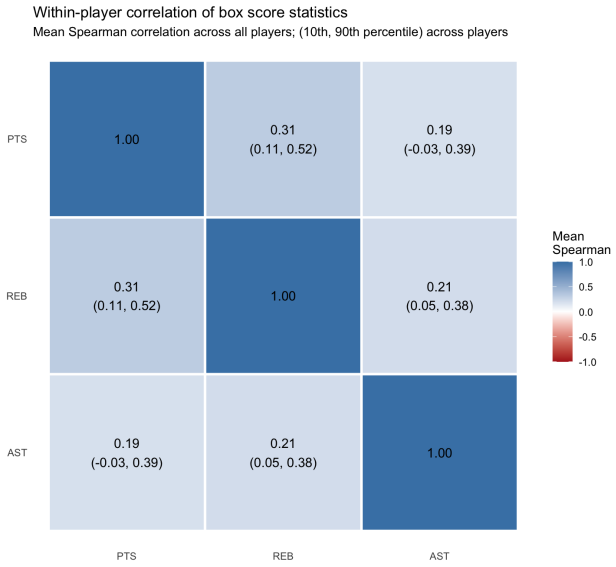


Figure 1: Within player correlation between Box Score Statistics

Copulas provide a natural framework for this problem, since they separate each variable's marginal distribution from the dependence structure that links them (Sklar, 1959). Once the marginals and the copula are estimated, we can simulate from the fitted joint distribution and estimate the probability of any combination of variable thresholds, including combinations with little or no historical support. Gaussian copulas were popularized in finance for modeling default correlation (Li, 2000), and copula methods have been applied to basketball data before; D'Urso et al. (2023) used a copula-based Bayesian network to study the determinants of player performance, and Liu and Yu (2026) used a copula-based Dirichlet process mixture to find subgroups of similar players. Our objective is to model the joint distribution of points, rebounds, and assists at the individual player level.

In this paper, we model this joint distribution using a Gaussian copula, fitting the model in a strictly out-of-sample setting by re-estimating it before every game. We evaluate joint exceedance forecasts across 27 combinations of player-relative thresholds with Log Loss as our loss function, comparing the Gaussian copula against independence and empirical-frequency benchmarks. The results show that the benefits of modeling dependence are concentrated in combinations where all three thresholds are high. The model uses only a player's own game history and does not account for opponents, minutes, or pace. It therefore provides a baseline for assessing the predictive value of dependence modeling rather than a complete model pipeline.

2 Gaussian Copulas

A copula is a multivariate distribution function on the unit hypercube $[0,1]^d$ whose univariate margins are uniform. Sklar's theorem (Sklar, 1959) states that for any d-dimensional joint distribution function F with marginals $F_1, \ldots, F_d$, there exists a copula C such that

$$F(x_1, \ldots, x_d) = C\big(F_1(x_1), \ldots, F_d(x_d)\big), \tag{1}$$

and C is unique when the margins are continuous. The marginals describe each variable by itself, and the copula carries the information about their joint behavior, meaning the two can be specified and estimated separately. When the margins are discrete, as with box score counts, C is not unique; we return to this in Section 4.

Let Φ denote the standard normal distribution function and Φ^{-1} its quantile function. The Gaussian copula is the dependence structure of a multivariate normal distribution,

$$C_{\mathbf{R}}(u_1, \ldots, u_d) = \Phi_{\mathbf{R}}\big(\Phi^{-1}(u_1), \ldots, \Phi^{-1}(u_d)\big), \tag{2}$$

where $\Phi_{\mathbf{R}}$ is the d-variate standard normal distribution function with correlation matrix $\mathbf{R}$. The correlation matrix fully determines the dependence. The Gaussian copula captures dependence through a single correlation matrix, allowing each pair of statistics its own strength of association while retaining flexibility in the marginals.

3 Estimation

For each player we model $\mathbf{X} = (X_{\text{pts}}, X_{\text{reb}}, X_{\text{ast}})$ with a Gaussian copula and empirical marginals. To mimic a genuine forecasting setting, the model is re-estimated before every game using only past games: for game $t \geq 21$, the training set consists of the $n = t - 1$ preceding games $\mathbf{x}_1, \ldots, \mathbf{x}_n$.

Marginals. For $j \in \{\text{pts}, \text{reb}, \text{ast}\}$, the marginal $\hat{F}_j$ is the empirical distribution function of the player's training games, with quantile function $\hat{Q}_j = \hat{F}_j^{-1}$ given by the linearly interpolated sample quantile.

Dependence. Each observation is mapped to the unit interval through its rank-based pseudo-observation and then to the normal scale,

$$\hat{u}_{sj} = \frac{r_{sj}}{n+1}, \qquad z_{sj} = \Phi^{-1}(\hat{u}_{sj}), \qquad s = 1, \ldots, n, \tag{3}$$

where r_{sj} is the rank of x_{sj} among the training games, with average ranks for ties. The copula's correlation matrix $\hat{R}$ is estimated as the sample correlation matrix of the normal scores $z_s = (z_{s,\text{pts}}, z_{s,\text{reb}}, z_{s,\text{ast}})$. We use the 'nearPD' procedure from the R package 'Matrix' (Bates et al., 2026) to obtain the nearest positive definite correlation matrix, enforcing a unit diagonal. The fitted joint distribution is

$$\hat{F}(x_{\text{pts}}, x_{\text{reb}}, x_{\text{ast}}) = \Phi_{\hat{\mathbf{R}}}\big(\Phi^{-1}(\hat{F}_{\text{pts}}(x_{\text{pts}})),\, \Phi^{-1}(\hat{F}_{\text{reb}}(x_{\text{reb}})),\, \Phi^{-1}(\hat{F}_{\text{ast}}(x_{\text{ast}}))\big). \tag{4}$$

Simulation. We estimate joint exceedance probabilities by Monte Carlo. For $b = 1, \ldots, B$ we draw

$$\mathbf{z}^{(b)} \sim N(\mathbf{0}, \hat{\mathbf{R}}), \qquad u_j^{(b)} = \Phi\big(z_j^{(b)}\big), \qquad x_j^{(b)} = \hat{Q}_j\big(u_j^{(b)}\big), \tag{5}$$

giving simulated games $\mathbf{x}^{(b)}$ with the player's own marginals and the estimated dependence. We use $B = 5{,}000$. Fitting and simulation used the `copula` package in R (Jun Yan, 2007).

4 Evaluation Design

4.1 Thresholds and events

For each player-game, we use the player's training games (at least 20) to set player-specific thresholds. For quartile levels $\boldsymbol{\alpha} = (\alpha_{\text{pts}}, \alpha_{\text{reb}}, \alpha_{\text{ast}})$ with each $\alpha_j \in \{0.25, 0.50, 0.75\}$, giving $3^3 = 27$ combinations, the thresholds are $\tau_j = \hat{Q}_j(\alpha_j)$ and the event of interest is

$$A_{\boldsymbol{\alpha}} = \{X_{\text{pts}} \geq \tau_{\text{pts}},\ X_{\text{reb}} \geq \tau_{\text{reb}},\ X_{\text{ast}} \geq \tau_{\text{ast}}\}. \tag{6}$$

For example, the combination (PTS Q50, REB Q25, AST Q75) is the event that the player scores at least his historical median in points, at least his 25th percentile in rebounds, and at least his 75th percentile in assists. Because thresholds are computed from each player's own history, a "high" threshold is relative to that player and the events are comparable across players with different statistical profiles. Thresholds are updated along with the model before every game. We score each forecast against the realized outcome $y_t = \mathbf{1}\{\mathbf{x}_t \in A_{\boldsymbol{\alpha}}\}$, where $\mathbf{x}_t$ is the stat line of the game being predicted, resulting in 27 binary outcomes for each player every game.

Although the Gaussian copula can theoretically estimate the likelihood of any combination of thresholds, we limit our analysis to quartiles (Q25, Q50, Q75) to limit the complexity of the prediction space and retain explanatory clarity.

Points, rebounds, and assists are integer-valued, so their marginal distributions are discrete with frequent ties. Two consequences follow. First, the copula is not unique, so $\hat{F}$ is one of many joint distributions consistent with the marginals. Second, because the event uses $\geq$ and the thresholds are interpolated quantiles, $\Pr(X_j \geq \tau_j)$ can differ from $1-\alpha_j$, typically being larger. Thresholds are not rounded to the nearest integer; exceedance events are evaluated directly using the interpolated thresholds. We use the same threshold definition for the copula and both benchmarks, ensuring a like-for-like comparison.

4.2 Forecasts and benchmarks

The copula forecast of $\Pr(A_{\boldsymbol{\alpha}})$ is the Monte Carlo estimate

$$\hat{p}^{\text{cop}} = \frac{1}{B} \sum_{b=1}^{B} \mathbf{1}\{\mathbf{x}^{(b)} \in A_{\boldsymbol{\alpha}}\}. \tag{7}$$

The *empirical* benchmark is the frequency of the event in the n training games,

$$\hat{p}^{\text{emp}} = \frac{1}{n} \sum_{s=1}^{n} \mathbf{1}\{\mathbf{x}_s \in A_{\boldsymbol{\alpha}}\}. \tag{8}$$

The *independence* benchmark multiplies the empirical marginal exceedance frequencies,

$$\hat{p}^{\text{ind}} = \prod_{j=1}^{3} \left[\frac{1}{n} \sum_{s=1}^{n} \mathbf{1}\{x_{sj} \geq \tau_j\} \right]. \tag{9}$$

This captures the marginal distributions but assumes independence.

4.3 Loss and inference

Forecasts are scored with log loss:

$$\ell(p, y) = -\big[y \log p + (1 - y) \log(1 - p)\big], \tag{10}$$

With sample sizes as small as 20, the empirical joint exceedance probability can be exactly zero. We therefore clip p to $[10^{-6}, 1 - 10^{-6}]$ so that forecasts of exactly zero or one have finite loss.

4.4 Data

We sampled one NBA starter from each team in the first game of each two season window, selecting players who appeared in that game and had at least 20 total games across the two season window. All player-game observations were restricted to games in which the player recorded non-zero minutes. Players were then followed across the two consecutive seasons, with forecasts estimated before each game. Players who were traded or missed games due to injury remained in the sample, provided they met the initial eligibility criteria. We used seven windows, each covering two consecutive seasons, and avoiding the 2020 COVID season. The windows include 2010–11, 2012–13, 2014–15, 2016–17, 2018–19, 2023–24, and 2025–26. This yielded 207 players and approximately 22,000 player-game binary outcomes.

5 Results

Figure 2 and 3 compare the copula with each benchmark on the same player-games. For each threshold combination, it shows the relative improvement, $1 - \ell_{\text{Gaussian}}/\ell_{\text{comparison}}$. Because both models forecast the same events, this is a paired comparison: player-games that are hard to forecast are hard for both models, and this shared difficulty largely cancels in the ratio. Confidence intervals come from a cluster bootstrap that resamples player-runs and recomputes both losses on each resampled set, so they respect dependence among games from the same player.

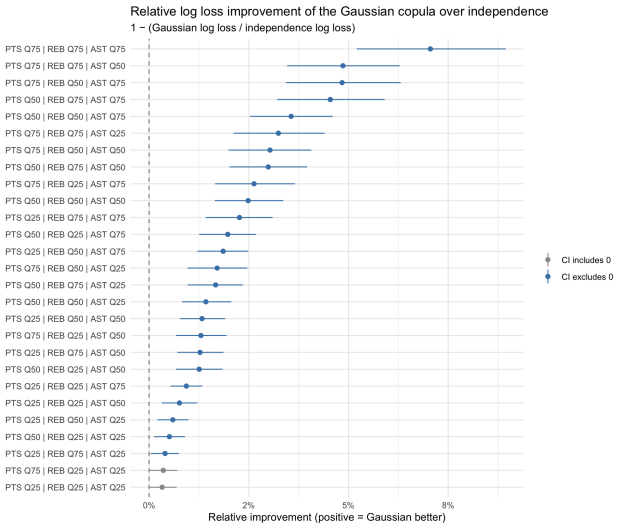


Figure 2: Relative log loss improvement of the Gaussian copula over the independence benchmark by threshold combination. Points show estimates and bars show confidence intervals.

We observe clear separation between the Gaussian copula and the independence model with all but two thresholds showing statistically significant improvements. For thresholds with multiple Q75s we see between $2\% - 8\%$ decreases in log loss, showing that the Gaussian copula is better at capturing the correlations at the upper quartiles.

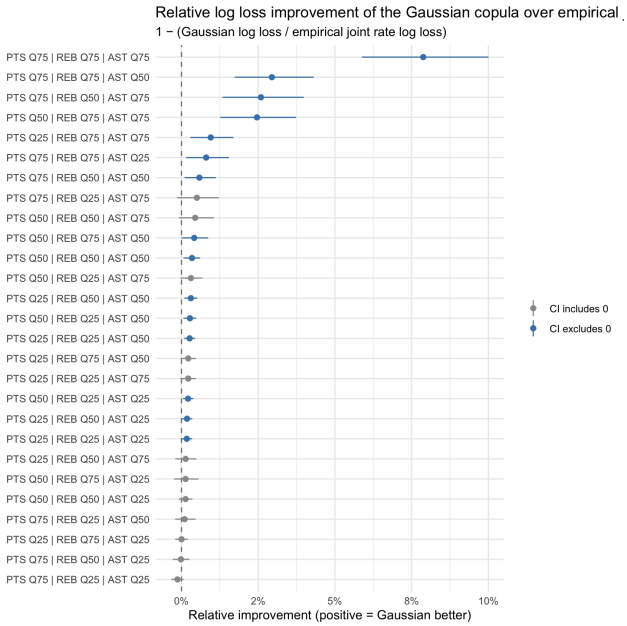


Figure 3: Relative log loss improvement of the Gaussian copula over the empirical benchmark by threshold combination. Points show estimates and bars show confidence intervals.

The log loss looks very similar for the Gaussian copula and empirical model for many of the lower thresholds. Only at the upper thresholds do we see the Gaussian copula performing significantly better. We attach in the Appendix 8 examples of Log Loss for more extreme thresholds which show even greater out performance.

For some context of the magnitude of Log Loss in Table 1 we show the average Log Loss across example thresholds.

Table 1: Average Log Loss Across Selected Threshold Combinations

Thresholds	Gaussian	Empirical	Independence
Q25, Q25, Q25	0.658	0.659	0.660
Q50, Q50, Q50	0.571	0.573	0.585
Q75, Q75, Q75	0.245	0.266	0.264

5.1 Calibration

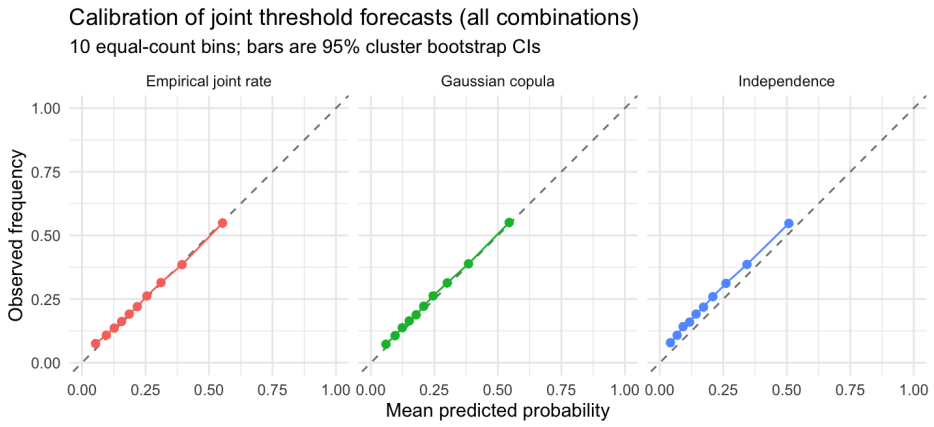
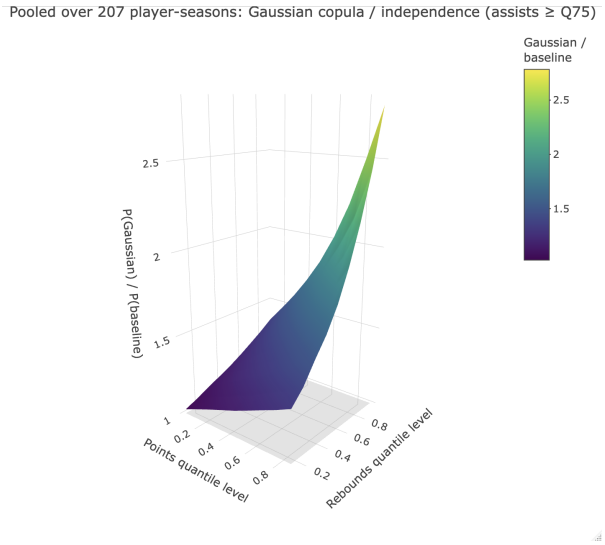


Figure 4: Calibration Plot for All Three Models

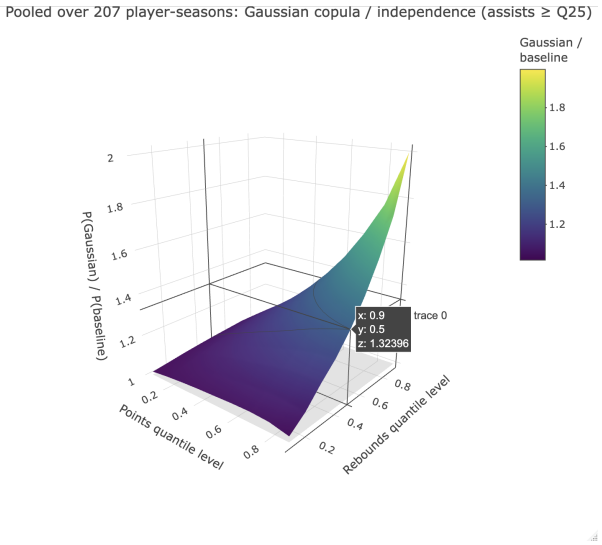
We observe that the empirical and Gaussian models are well calibrated, while the model assuming marginal independence systematically under predicts the probabilities of joint exceedance.

5.2 Relative Density Plots

Figure 5a shows the ratio of the copula's joint exceedance probability to the independence benchmark, fixing assists at the Q75 threshold and varying the points and rebounds quartile levels. Figure 5b is the same graph but with assists fixed at Q25. We see in both cases the ratio of the probabilities rising as the quartile level increases, with a more extreme increase when assists are conditioned at Q75.



(a) Independence benchmark comparison - Assists at Q75



(b) Independence benchmark comparison - Assists at Q25

Figure 5: Ratio of the Gaussian copula's joint exceedance probability relative to the independence benchmark.

Although the Gaussian distribution places relatively little probability mass in its tails, we observe that this does not prevent the Gaussian copula from capturing dependence at high quantiles. The copula separates the dependence structure from the marginal distributions, which in our case are empirical and therefore retain the observed tail behavior of each statistic.

6 Conclusion

A Gaussian copula performs on par or better than the empirical joint frequency model at lower threshold levels, while demonstrating clear out performance over both the independence benchmark and the empirical model at higher quartiles. The result is a simple, a fully generative model that can be used to calculate threshold probabilities of both observed and previously unobserved stat lines. It represents a baseline that future models could extend with richer pre-game context.

7 Discussion and Future Directions

The results support modeling dependence: gains over independence are consistent across all 27 combinations, and gains over the empirical estimate are largest where joint events are rare and the empirical frequency is noisiest. Several limitations remain. The Gaussian copula imposes a symmetric dependence structure and exhibits zero asymptotic tail dependence. We explored alternative dependence specifications, including the Gumbel copula, as well as approaches that did not impose a parametric joint distribution and allowed the dependence structure to vary across players. We didn't find any clear improvement in predictive performance, so we retained the Gaussian copula as our primary model specification.

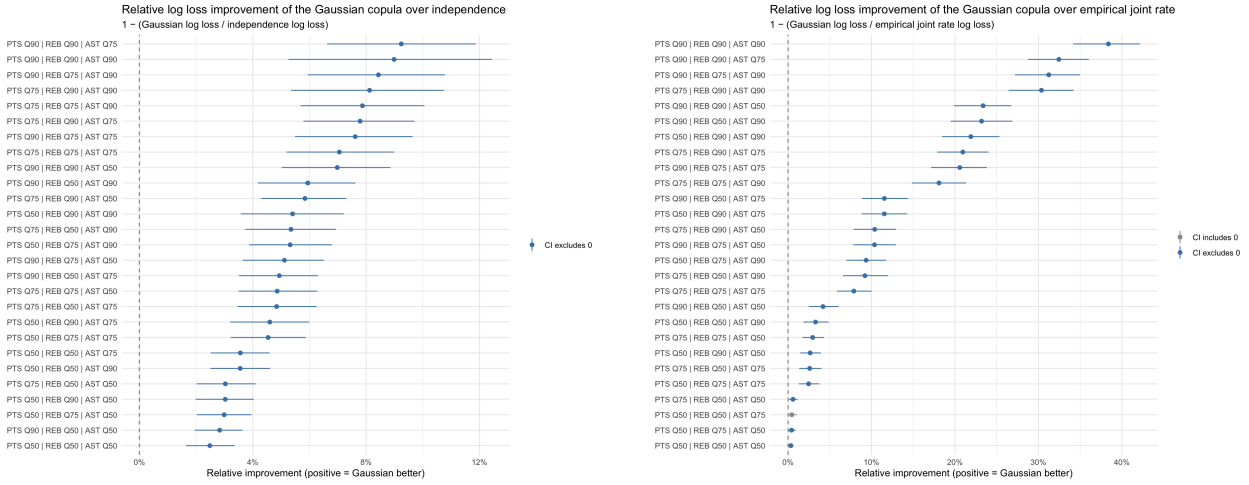
Applications using Gaussian copulas such as (Li, 2000) have been criticized for their lack of sensitivity to rare events and tail risk. However, this isn't particularly relevant in this situation since we're predicting binary outcomes, i.e. there aren't extreme asymmetric gains/losses to an incorrect prediction of an outlier.

Our model uses only a player's own history, so it ignores ex ante knowledge of opponents, minutes, and pace. We suspect that conditioning on pre-game predicted possessions would improve the joint estimates, since playing time is the main driver of the dependence we model. Finally, we evaluate player-relative quartile thresholds rather than the absolute lines offered by sports books, so we have not tested out performance against market implied odds.

8 Appendix

All code and plots are accessible on GitHub.

We attach graphs of relative Log Loss improvement for more extreme thresholds.



(a) Gaussian vs Independence

(b) Gaussian vs Empirical

Figure 6: Relative Log Loss Improvement at High Thresholds

Table 2 and 3 show the mean paired differences between the Gaussian copula and alternative method, alongside bootstrapped p-values.

Table 2: Independence vs. Gaussian Copula

Combo	Mean Difference	p_{boot}
PTS Q75 — REB Q75 — AST Q75	0.01860	0.0001
PTS Q75 — REB Q50 — AST Q75	0.01670	0.0001
PTS Q75 — REB Q75 — AST Q50	0.01660	0.0001
PTS Q50 — REB Q50 — AST Q75	0.01650	0.0001
PTS Q50 — REB Q75 — AST Q75	0.01580	0.0001
PTS Q50 — REB Q50 — AST Q50	0.01450	0.0001
PTS Q50 — REB Q75 — AST Q50	0.01350	0.0001
PTS Q75 — REB Q50 — AST Q50	0.01350	0.0001
PTS Q75 — REB Q75 — AST Q25	0.01230	0.0001
PTS Q50 — REB Q25 — AST Q75	0.01020	0.0001
PTS Q75 — REB Q25 — AST Q75	0.01000	0.0001
PTS Q25 — REB Q50 — AST Q75	0.00977	0.0001
PTS Q50 — REB Q50 — AST Q25	0.00906	0.0001
PTS Q25 — REB Q75 — AST Q75	0.00893	0.0001
PTS Q25 — REB Q50 — AST Q50	0.00852	0.0001
PTS Q75 — REB Q50 — AST Q25	0.00847	0.0001
PTS Q50 — REB Q75 — AST Q25	0.00835	0.0001
PTS Q50 — REB Q25 — AST Q50	0.00798	0.0001
PTS Q25 — REB Q75 — AST Q50	0.00651	0.0001
PTS Q75 — REB Q25 — AST Q50	0.00639	0.0001
PTS Q25 — REB Q25 — AST Q75	0.00545	0.0001
PTS Q25 — REB Q25 — AST Q50	0.00509	0.0012
PTS Q25 — REB Q50 — AST Q25	0.00403	0.0032
PTS Q50 — REB Q25 — AST Q25	0.00345	0.0104
PTS Q25 — REB Q75 — AST Q25	0.00225	0.0248
PTS Q25 — REB Q25 — AST Q25	0.00217	0.0734
PTS Q75 — REB Q25 — AST Q25	0.00195	0.0480

Table 3: Empirical vs. Gaussian Copula

Combo	Mean Difference	p_{boot}
PTS Q75 — REB Q75 — AST Q75	0.02100	0.0001
PTS Q75 — REB Q75 — AST Q50	0.00989	0.0001
PTS Q75 — REB Q50 — AST Q75	0.00872	0.0001
PTS Q50 — REB Q75 — AST Q75	0.00839	0.0001
PTS Q25 — REB Q75 — AST Q75	0.00371	0.0038
PTS Q75 — REB Q75 — AST Q25	0.00298	0.0124
PTS Q75 — REB Q50 — AST Q50	0.00253	0.0150
PTS Q50 — REB Q50 — AST Q75	0.00202	0.1090
PTS Q50 — REB Q50 — AST Q50	0.00196	0.0150
PTS Q25 — REB Q50 — AST Q50	0.00194	0.0028
PTS Q75 — REB Q25 — AST Q75	0.00187	0.1240
PTS Q50 — REB Q75 — AST Q50	0.00184	0.0422
PTS Q25 — REB Q25 — AST Q50	0.00178	0.0042
PTS Q50 — REB Q25 — AST Q50	0.00174	0.0064
PTS Q50 — REB Q25 — AST Q75	0.00156	0.0746
PTS Q50 — REB Q25 — AST Q25	0.00145	0.0098
PTS Q25 — REB Q25 — AST Q75	0.00128	0.0830
PTS Q25 — REB Q50 — AST Q25	0.00122	0.0326
PTS Q25 — REB Q25 — AST Q25	0.00115	0.0592
PTS Q25 — REB Q75 — AST Q50	0.00112	0.0818
PTS Q50 — REB Q50 — AST Q25	0.000845	0.2550
PTS Q25 — REB Q50 — AST Q75	0.000704	0.4330
PTS Q50 — REB Q75 — AST Q25	0.000664	0.5250
PTS Q75 — REB Q25 — AST Q50	0.000507	0.5580
PTS Q25 — REB Q75 — AST Q25	0.000015	0.9820
PTS Q75 — REB Q50 — AST Q25	-0.000079	0.9140
PTS Q75 — REB Q25 — AST Q25	-0.000705	0.2090

References

Bates, D., Maechler, M., and Jagan, M. (2026). *Matrix: Sparse and Dense Matrix Classes and Methods*. R package version 1.7-5.

D'Urso, P., De Giovanni, L., and Vitale, V. (2023). A Bayesian network to analyse basketball players' performances: a multivariate copula-based approach. *Annals of Operations Research*, 325:419–440.

Jun Yan (2007). Enjoy the joy of copulas: With a package copula. *Journal of Statistical Software*, 21(4):1–21.

Li, D. X. (2000). On default correlation: A copula function approach. Technical Report 99-07, The RiskMetrics Group.

Liu, Y. and Yu, S. (2026). Bayesian Dirichlet process copula mixtures for heterogeneous multi-

cluster data: Methods and an NBA player stats application. *Statistical Analysis and Data Mining: An ASA Data Science Journal*, 19(3):e70074.

Sklar, A. (1959). Fonctions de répartition à n dimensions et leurs marges. *Publications de l'Institut de Statistique de l'Université de Paris*, 8:229–231. Distribution functions with n dimensions and their margins.

Whelan, K. (2026). The parlay puzzle: Expected utility and multi-leg betting. UCD Centre for Economic Research Working Paper Series WP2026/10, University College Dublin, School of Economics.